

Chapter 8

Loglinear and Logit Models for Contingency Tables

{ch:loglin}

Loglinear models comprise another special case of generalized linear models designed for contingency tables of frequencies. They are most easily interpreted through visualizations, including mosaic displays and plots of associated logit models. As with logistic regression, diagnostic plots and influence plots help to assure that the fitted model is an adequate summary of associations among variables.

8.1 Introduction

{sec:loglin-intro}

Tables are like cobwebs, like the sieve of Danaides; beautifully reticulated, orderly to look upon, but which will hold no conclusion. Tables are abstractions, and the object a most concrete one, so difficult to read the essence of.

From *Chartism* by Thomas Carlyle (1840), Chapter II, Statistics

The chapter continues the modeling framework begun in Chapter 7, and takes up the case of loglinear models for contingency tables of frequencies, when all variables are discrete, another special case of generalized linear models. These models provide a comprehensive scheme to describe and understand the associations among two or more categorical variables. Whereas logistic regression models focus on the prediction of one response factor, loglinear models treat all variables symmetrically, and attempt to model all important associations among them.

In this sense, loglinear models are analogous to a correlation analysis of continuous variables, where the goal is to determine the patterns of dependence and independence among a set of variables. Such models are also particularly useful when there are two or more response variables, a case that would require a multivariate version of the generalized linear model, for which the current theory and implementations are thin at best. Nonetheless, when one variable is indeed a response and the others are explanatory, certain loglinear models are equivalent to logistic models for that response.

Chapter 5 and Chapter 6 introduced some basic aspects of loglinear models in connection with mosaic displays and correspondence analysis. In this chapter, the focus is on fitting and interpreting loglinear models. The usual analyses, with `loglm()` and `glm()` present the results

in terms of tables of parameter estimates. Particularly for larger tables, it becomes difficult to understand the nature of these associations from tables of parameter estimates. Instead, we emphasize plots of observed and predicted frequencies, probabilities or log odds (when there are one or more response variables), as well as mosaic and other displays for interpreting a given model, and residual and influence plots for model diagnostics. We also illustrate how mosaic displays and correspondence analysis plots may be used in a complementary way to the usual numerical summaries, to provide additional insights into the data.

Section 8.2 gives a brief overview of loglinear models in relation to the more familiar ANOVA and regression models for quantitative data. Methods and software for fitting these models are discussed in Section 8.3. When one variable is a response, logit models for that response provide a simpler, but equivalent means for interpreting and graphing results of loglinear models, as we describe in Section 8.4. Another class of simplified models (Section 8.6) occurs when one or more of the explanatory variables are ordinal, and discrete levels might be replaced by numerical values.

TODO: Complete Chapter overview of sections

8.2 Loglinear models for frequencies

Loglinear models have been developed from two formally distinct, but related perspectives. The first is a discrete analog of familiar ANOVA models for quantitative data, where the multiplicative relations among joint and marginal probabilities are transformed into an additive one by transforming the counts to logarithms. The second is an analog of regression models, where the log of the cell frequency is modeled as a linear function of discrete predictors, with a random component often taken as the Poisson distribution and called *Poisson regression*.

8.2.1 Loglinear models as ANOVA models for frequencies

For two discrete variables, A and B , suppose we have a multinomial sample of n_{ij} observations in each cell i, j of an $I \times J$ contingency table. To ease notation, we replace a subscript by $+$ to represent summation over that dimension, so that $n_{i+} = \sum_j n_{ij}$, $n_{+j} = \sum_i n_{ij}$, and $n_{++} = \sum_{ij} n_{ij}$.

Let π_{ij} be the joint probabilities in the table, and let $m_{ij} = n_{++}\pi_{ij}$ be the expected cell frequencies under any model. Conditional on the observed total count, n_{++} , each count has a Poisson distribution, with mean m_{ij} . Any loglinear model may be expressed as a linear model for the log m_{ij} . For example, the hypothesis of independence means that the expected frequencies, m_{ij} , obey

$$m_{ij} = \frac{m_{i+} m_{+j}}{m_{++}} .$$

This multiplicative model can be transformed to an additive (linear) model by taking logarithms of both sides:

$$\log(m_{ij}) = \log(m_{i+}) + \log(m_{+j}) - \log(m_{++}) ,$$

which is usually expressed in an equivalent form in terms of model parameters,

$$\log(m_{ij}) = \mu + \lambda_i^A + \lambda_j^B \tag{8.1}$$

where μ is a function of the total sample size, λ_i^A is the “main effect” for variable A, $\lambda_i^A = \log \pi_{i+} - \sum_k (\log \pi_{k+})/I$, and λ_j^B is the “main effect” for variable B, $\lambda_j^B = \log \pi_{+j} - \sum_k (\log \pi_{+k})/J$. Model Eqn. (8.1) is called the **loglinear independence model** for a two-way table.

In this model, there are $1+I+J$ parameters, but only $(I-1)+(J-1)$ are separately estimable. Hence, the typical ANOVA sum-to-zero restrictions are usually applied to the parameters:

$$\sum_i \lambda_i^A = \sum_j \lambda_j^B = 0 .$$

These “main effects” in loglinear models pertain to differences among the marginal probabilities of a variable (which are usually not of direct interest).

Other restrictions to make the parameters identifiable are also used. Setting the first values, λ_1^A and λ_1^B to zero (the default in `glm()`), defines $\lambda_i^A = \log \pi_{i+} - \log \pi_{1+}$, and $\lambda_j^B = \log \pi_{+j} - \log \pi_{+1}$, as deviations from the first, reference category, but these parameterizations are otherwise identical. For modeling functions in `R` (`lm()`, `glm()`, etc.) the reference category parameterization is obtained using `contr.treatment()`, while the sum-to-zero constraints are obtained with `contr.sum()`.

Model Eqn. (8.1) asserts that the row and column variables are independent. For a two-way table, a model that allows an arbitrary association between the variables is the **saturated model**, including an additional term, λ_{ij}^{AB} :

$$\log(m_{ij}) = \mu + \lambda_i^A + \lambda_j^B + \lambda_{ij}^{AB} , \quad (8.2) \quad \{\text{eq:lsat}\}$$

where again, restrictions must be imposed for estimation:

$$\sum_i \lambda_i^A = 0, \quad \sum_j \lambda_j^B = 0, \quad \sum_i \lambda_{ij}^{AB} = \sum_j \lambda_{ij}^{AB} = 0 . \quad (8.3) \quad \{\text{eq:lrestrict}\}$$

There are thus $I-1$ linearly independent λ_i^A row parameters, $J-1$ linearly independent λ_j^B column parameters, and $(I-1)(J-1)$ linearly independent λ_{ij}^{AB} association parameters. This model is called the **saturated model** because the number of parameters in μ , λ_i^A , λ_j^B , and λ_{ij}^{AB} is equal to the number of frequencies in the two-way table,

$$\underbrace{1}_{(\mu)} + \underbrace{I-1}_{(\lambda_i^A)} + \underbrace{J-1}_{(\lambda_j^B)} + \underbrace{(I-1)(J-1)}_{(\lambda_{ij}^{AB})} = \underbrace{IJ}_{(n_{ij})}$$

The association parameters λ_{ij}^{AB} express the departures from independence, so large absolute values pertain to cells that differ from the independence model.

Except for the difference in notation, model Eqn. (8.2) is formally the same as a two-factor ANOVA model with an interaction, typically expressed as $E(y_{ij}) = \mu + \alpha_i + \beta_j + (\alpha\beta)_{ij}$. Hence, associations between variables in loglinear models are analogous to interactions in ANOVA models. The use of superscripted symbols, λ_i^A , λ_j^B , λ_{ij}^{AB} rather than separate Greek letters is a convention in loglinear models, and useful mainly for multiway tables.

Models such as Eqn. (8.1) and Eqn. (8.2) are examples of **hierarchical models**. This means that the model must contain all lower-order terms contained within any high-order term in the model. Thus, the saturated model, Eqn. (8.2) contains λ_{ij}^{AB} , and therefore *must* contain λ_i^A and λ_j^B . As a result, hierarchical models may be identified by the shorthand notation which lists only the high-order terms: model Eqn. (8.2) is denoted $[AB]$, while model Eqn. (8.1) is $[A][B]$.

8.2.2 Loglinear models for three-way tables

{sec:loglin

Loglinear models for three-way contingency tables were described briefly in Section 5.4.1. Each type of model allows associations among different sets of variables and each has a different independence interpretation, as illustrated in Table 5.2.

For a three-way table, the saturated model, denoted $[ABC]$ is

$$\{eq:lsat3\} \quad \log m_{ijk} = \mu + \lambda_i^A + \lambda_j^B + \lambda_k^C + \lambda_{ij}^{AB} + \lambda_{ik}^{AC} + \lambda_{jk}^{BC} + \lambda_{ijk}^{ABC} . \quad (8.4)$$

This model allows all variables to be associated; Eqn. (8.4) fits the data perfectly because the number of independent parameters equals the number of table cells. Two-way terms, such as λ_{ij}^{AB} pertain to the *conditional association* between pairs of factors, controlling for the remaining variable. The presence of the three-way term, λ_{ijk}^{ABC} , means that the partial association (conditional odds ratio) between any pair varies over the levels of the third variable.

Omitting the three-way term in Model Eqn. (8.4) gives the model $[AB][AC][BC]$,

$$\{eq:lno3way\} \quad \log m_{ijk} = \mu + \lambda_i^A + \lambda_j^B + \lambda_k^C + \lambda_{ij}^{AB} + \lambda_{ik}^{AC} + \lambda_{jk}^{BC} , \quad (8.5)$$

in which all pairs are conditionally dependent given the remaining one. For any pair, the conditional odds ratios are the *same* at all levels of the remaining variable, so this model is often called the **homogeneous association model**.

The interpretation of terms in this model may be illustrated using the Berkeley admissions data (Example 4.10 and Example 4.14), for which the factors are Admit, Gender, and Department, in a $2 \times 2 \times 6$ table. In the homogeneous association model,

$$\{eq:berk1\} \quad \log m_{ijk} = \mu + \lambda_i^A + \lambda_j^D + \lambda_k^G + \lambda_{ij}^{AD} + \lambda_{ik}^{AG} + \lambda_{jk}^{DG} , \quad (8.6)$$

the λ -parameters have the following interpretations:

- The main effects, λ_i^A , λ_j^D and λ_k^G pertain to differences in the one-way marginal probabilities. Thus λ_j^D relates to differences in the total number of applicants to these departments, while λ_k^G relates to the differences in the overall numbers of men and women applicants.
- λ_{ij}^{AD} describes the conditional association between admission and department, that is different admission rates across departments (controlling for gender).
- λ_{ik}^{AG} relates to the conditional association between admission and gender, controlling for department. This term, if significant, might be interpreted as indicating gender-bias in admissions.
- λ_{jk}^{DG} , the association between department and gender, indicates whether males and females apply differentially across departments.

As we discussed earlier (Section 5.4), loglinear models for three-way (and larger) tables often have an interpretation in terms of various types of independence relations illustrated in Table 5.2. The model Eqn. (8.5) has no such interpretation, however the smaller model $[AC][BC]$ can be interpreted as asserting that A and B are (conditionally) independent controlling for C ; this independence interpretation is symbolized as $A \perp B | C$. Similarly, the model $[AB][C]$ asserts that A and B are jointly independent of C : $(A, B) \perp C$, while the model $[A][B][C]$ is the model of mutual (complete) independence, $A \perp B \perp C$.

8.2.3 Loglinear models as GLMs for frequencies

In the GLM approach, a loglinear model may be cast in the form of a regression model for $\log \mathbf{m}$, where the table cells are reshaped to a column vector. One advantage is that models for tables of any size and structure may be expressed in a compact form.

For a contingency table of variables A, B, C, \dots , with $N = I \times J \times K \times \dots$ cells, let \mathbf{n} denote a column vector of the observed counts arranged in standard order, and let \mathbf{m} denote a similar vector of the expected frequencies under some model. Then *any* loglinear model may be expressed in the form

$$\log \mathbf{m} = \mathbf{X}\boldsymbol{\beta} ,$$

where \mathbf{X} is a known design or **model matrix** and $\boldsymbol{\beta}$ is a column vector containing the unknown λ parameters.

For example, for a 2×2 table, the saturated model Eqn. (8.2) with the usual zero-sum constraints Eqn. (8.3) can be represented as

$$\log \begin{pmatrix} m_{11} \\ m_{12} \\ m_{21} \\ m_{22} \end{pmatrix} = \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 1 & -1 & -1 \\ 1 & -1 & 1 & -1 \\ 1 & -1 & -1 & 1 \end{bmatrix} \begin{pmatrix} \mu \\ \lambda_1^A \\ \lambda_1^B \\ \lambda_{11}^{AB} \end{pmatrix}$$

Note that only the linearly independent parameters are represented here. $\lambda_2^A = -\lambda_1^A$, because $\lambda_1^A + \lambda_2^A = 0$, and $\lambda_2^B = -\lambda_1^B$, because $\lambda_1^B + \lambda_2^B = 0$, and so forth.

An additional substantial advantage of the GLM formulation is that it makes it easier to express models with ordinal or quantitative variables. `glm()`, with a model formula of the form `Freq ~ .` involving factors A, B, \dots and quantitative variables x_1, x_2, \dots , constructs the model matrix \mathbf{X} from the terms given in the formula. A factor with K levels gives rise to $K - 1$ columns for its main effect and sets of $K - 1$ columns in each interaction effect. A quantitative predictor, say x_1 (with a linear effect) creates a single column with its values and interactions with other terms are calculated at the products of the columns for the main effects.

The parameterization for factors is controlled by the contrasts assigned to a given factor (if any), or by the general `contrasts` option, that gives the contrast functions used for unordered and ordered factors:

```
options("contrasts")

## $contrasts
##          unordered          ordered
## "contr.treatment" "contr.poly"
```

This says that, by default, unordered factors use the baseline (first) reference-level parameterization, while ordered factors are given a parameterization based on orthogonal polynomials, allowing linear, quadratic, ... effects, assuming integer-spacing of the factor levels.

8.3 Fitting and testing loglinear models

For a given table, possible loglinear models range from the baseline model of mutual independence, $[A][B][C][\dots]$ to the saturated model, $[ABC\dots]$ that fits the observed frequencies perfectly, but offers no simpler description or interpretation than the data itself.

{sec:loglin-fitting}

Fitting a loglinear model is usually a process of deciding which association terms are large enough (“significantly different from zero”) to warrant inclusion in a model to explain the observed frequencies. Terms which are excluded from the model go into the residual or error term, which reflects the overall badness-of-fit of the model. The usual goal of loglinear modeling is to find a small model (few association terms) which nonetheless achieves a reasonable fit (small residuals).

8.3.1 Model fitting functions

{sec:loglin-functions}

In R, the most basic function for fitting loglinear models is `loglin()` in the **stats** package. This uses the classical iterative proportional fitting (IPF) algorithm described in Haberman (1972) and Fienberg (1980, §3.4). It is designed to work with the frequency data in table form, and a model specified in terms of the (high-order) table margins to be fitted. For example, the model Eqn. (8.5) of homogeneous association for a three-way table is specified as

```
loglin(mytable, margin=list(c(1, 2), c(1, 3), c(2, 3)))
```

The function `loglm()` in **MASS** provides a more convenient front-end to `loglin()` to allow loglinear models to be specified using a model formula. With table variables A, B and C, the same model can be fit using `loglm()` as

```
loglm(~ (A + B + C)^2, data=mytable)
```

When the data is a frequency data frame with frequencies in `Freq`, for example, the result of `mydf <- as.data.frame(mytable)`, you can also use a two-sided formula:

```
loglm(Freq ~ (A + B + C)^2, data=mydf)
```

As implied in Section 8.2.3, loglinear models can also be fit using `glm()`, using `family=poisson` which constructs the model for `log(Freq)`. The same model is fit with `glm()` as:

```
glm(Freq ~ (A + B + C)^2, data=mydf, family=poisson)
```

While all of these fit equivalent models, the details of the printed output, model objects, and available methods differ, as indicated in some of the examples that follow.

It should be noted that both the `loglin()/loglm()` methods based on iterative proportional fitting, and the `glm()` approach using the Poisson model for log frequency give maximum likelihood estimates, $\hat{\mathbf{m}}$, of the expected frequencies, as long as all observed frequencies \mathbf{n} are *all* positive. Some special considerations when there cells with zero frequencies are described in Section ??.

8.3.2 Goodness-of-fit tests

{sec:loglin-goodfit}

For an n -way table, global goodness-of-fit tests for a loglinear model attempt to answer the question “How well does the model reproduce the observed frequencies?” That is, how close are the fitted frequencies estimated under the model to those of the saturated model or the data?

To avoid multiple subscripts for an n -way table, let $\mathbf{n} = n_1, n_2, \dots, n_N$ denote the observed frequencies in a table with N cells, and corresponding fitted frequencies $\hat{\mathbf{m}} = \hat{m}_1, \hat{m}_2, \dots, \hat{m}_N$

according to a particular loglinear model. The standard goodness-of-fit statistics are sums over the cells of measures of the difference between the \mathbf{n} and $\widehat{\mathbf{m}}$.

The most commonly used are the familiar Pearson chi-square,

{eq:pchi}

$$\chi^2 = \sum_i^N \frac{(n_i - \widehat{m}_i)^2}{\widehat{m}_i} , \quad (8.7)$$

and the likelihood-ratio G^2 or *deviance* statistic,

$$G^2 = 2 \sum_i^N n_i \log \left(\frac{n_i}{\widehat{m}_i} \right) . \quad (8.8) \quad \text{{eq:pgsq}}$$

Both of these statistics have asymptotic χ^2 distributions (as $\Sigma \mathbf{n} \rightarrow \infty$), reasonably well-approximated when all expected frequencies are large.¹ The (residual) degrees of freedom are the number of cells (N) minus the number of estimated parameters. The likelihood-ratio test can also be expressed as minus twice the difference in log-likelihoods under fitted and saturated models,

$$G^2 = -2 \log \left[\frac{\mathcal{L}(\widehat{\mathbf{m}}; \mathbf{n})}{\mathcal{L}(\mathbf{n}; \mathbf{n})} \right] = -2 [\log \mathcal{L}(\widehat{\mathbf{m}}; \mathbf{n}) - \log \mathcal{L}(\mathbf{n}; \mathbf{n})] ,$$

where $\mathcal{L}(\widehat{\mathbf{m}}; \mathbf{n})$ is the maximized likelihood for the fitted model and $\mathcal{L}(\mathbf{n}; \mathbf{n})$ is the corresponding likelihood for the saturated model.

In practice such global tests are less useful for comparing competing models. You may find that several different models have an acceptable fit or, sadly, that none do (usually because you are “blessed” with a large sample size). It is then helpful to compare competing models *directly*, and two strategies are particularly useful in these cases.

First, the likelihood-ratio G^2 statistic has the property in that one can compare two **nested models** by their difference in G^2 statistics, which has a χ^2 distribution on the difference in degrees of freedom. Two models, M_1 and M_2 , are nested when one, say, M_2 , is a special case of the other. That is, model M_2 (with ν_2 residual df) contains a subset of the parameters of M_1 (with ν_1 residual df), the remaining ones being effectively set to zero. Model M_2 is therefore more restrictive and cannot fit the data better than the more general model M_1 , i.e., $G^2(M_2) \geq G^2(M_1)$. The least restrictive of all models, with $G^2 = 0$ and $\nu = 0$ df is the saturated model for which $\widehat{\mathbf{m}} = \mathbf{n}$.

Assuming that the less restrictive model M_1 fits, the difference in G^2 ,

$$\Delta G^2 \equiv G^2(M_2 | M_1) = G^2(M_2) - G^2(M_1) \quad (8.9) \quad \text{{eq:gsqnest1}}$$

$$= 2 \sum_i n_i \log(\widehat{m}_{i1} / \widehat{m}_{i2}) \quad (8.10) \quad \text{{eq:gsqnest2}}$$

has a chi-squared distribution with $\text{df} = \nu_2 - \nu_1$. The last equality Eqn. (8.10) follows from substituting in Eqn. (8.8).

Rearranging terms in Eqn. (8.9), we see that we can partition the $G^2(M_2)$ into two terms,

$$G^2(M_2) = G^2(M_1) + G^2(M_2 | M_1) .$$

¹Except in bizarre or borderline cases, these tests provide the same conclusions when expected frequencies are at least moderate (all $\widehat{\mathbf{m}} > 5$). However, G^2 approaches the theoretical chi-squared distribution more slowly than does χ^2 , and the approximation may be poor when the average cell frequency is less than 5.

The first term measures the difference between the data and the more general model M_1 . If this model fits, the second term measures the additional lack of fit imposed by the more restrictive model. In addition to providing a more focused test, $G^2(M_2 | M_1)$ also follows the chi-squared distribution more closely when some $\{m_i\}$ are small (Agresti, 2013, §10.6.3).

Alternatively, a second strategy uses other measures that combine goodness-of-fit with model parsimony and may also be used to compare non-nested models. The statistics described below are all cast in the form of badness-of-fit relative to degrees of freedom, so that smaller values reflect “better” models.

The simplest idea (Goodman, 1971) is to use G^2/df (or χ^2/df), which has an asymptotic expected value of 1 for a good-fitting model. This type of measure is not routinely reported by R software, but is easy to calculate from output.

The **Akaike Information Criterion** (AIC) statistic (Akaike, 1973) is a very general criterion for model selection with maximum likelihood estimation, based on the idea of maximizing the information provided by a fitted model. AIC is defined generally as

$$\text{AIC} = -2 \log \mathcal{L} + 2k$$

where $\log \mathcal{L}$ is the maximized log likelihood and k is the number of parameters estimated in the model. Better models correspond to *smaller* AIC. For loglinear models, minimizing AIC is equivalent to minimizing

$$\text{AIC}^* = G^2 - 2\nu ,$$

where ν is the residual df, but the values of AIC and AIC* differ by an arbitrary constant. This form is easier to calculate by hand from the output of any modeling function if AIC is not reported, or an `AIC()` method is not available.

A third statistic of this type is the **Bayesian Information Criterion** (BIC) due to Schwartz (1978) and Raftery (1986),

$$\text{BIC} = G^2 - \nu \log(n) ,$$

where n is the total sample size. Both AIC and BIC penalize the fit statistic for increasing number of parameters. BIC also penalizes the fit directly with sample size, and so expresses a preference for less complex models than AIC as the sample size increases.

8.3.3 Residuals for loglinear models

Test statistics such as G^2 can determine whether a model has significant lack of fit, and model comparison tests using $\Delta G^2 = G^2(M_2 | M_1)$ can assess whether the extra term(s) in model M_1 significantly improves the model fit. Beyond these tests, the pattern of residuals for individual cells offers important clues regarding the nature of lack of fit and can help suggests associations that could be accounted for better.

As with logistic regression models (Section 7.4.1), several types of residuals are available for loglinear models. For cell i in the vector form of the contingency table, the **raw residual** is simply the difference between the observed and fitted frequencies, $e_i = n_i - \hat{m}_i$.

The **Pearson residual** is the square root of the contribution of the cell to the Pearson χ^2 ,

$$r_i = \frac{n_i - \hat{m}_i}{\sqrt{\hat{m}_i}} \quad (8.11)$$

{eq:reschi2}

sec:loglin-residuals}

Similarly, the *deviance residual* can be defined as

$$g_i = \text{sign}(n_i - \hat{m}_i) \sqrt{2n_i \log(n_i / \hat{m}_i) - 2(n_i - \hat{m}_i)} \quad (8.12)$$

Both of these attempt to standardize the distribution of the residuals to a standard normal, $N(0, 1)$ form. However, as pointed out by Haberman (1973), the asymptotic variance of these is less than one (with average value df/N) but, worse—the variance decreases with \hat{m}_i . That is, residuals for cells with small expected frequencies have larger sampling variance, as might be expected.

Consequently, Haberman suggested dividing the Pearson residual by its estimated standard error, giving what are often called *adjusted residuals*. When loglinear models are fit using the GLM approach, the adjustment may be calculated using the leverage (“hat value”), h_i to give appropriately standardized residuals,

$$\begin{aligned} r_i^* &= r_i / \sqrt{1 - h_i} \\ g_i^* &= g_i / \sqrt{1 - h_i} \end{aligned}$$

These standardized versions are generally preferable, particularly for visualizing model lack of fit using mosaic displays. The reason for preferring adjusted residuals is illustrated in Figure 8.1, a plot of the factors, $\sqrt{1 - h_i}$, determining the standard errors of the residuals against the fitted values, \hat{m}_i , in the model for the UCBA admissions data described in Example 8.2 below. The values shown in this plot are calculated as:

```
berkeley <- as.data.frame(UCBA admissions)
berk.glm1 <- glm(Freq ~ Dept * (Gender+Admit), data=berkeley, family="poisson")
fit <- fitted(berk.glm1)
hat <- hatvalues(berk.glm1)
stderr <- sqrt(1-hat)
```

In R, raw, Pearson and deviance residuals may be obtained using `residuals(model, type=)`, where `type` is one of "raw", "pearson" and "deviance". Standardized (adjusted) residuals can be calculated using `rstandard(model, type=)`, for `type="pearson"` and `type="deviance"` versions.

8.3.4 Using `loglm()`

{loglin-loglin}

Here we illustrate the basics of fitting loglinear models using `loglm()`. As indicated in Section 8.3.1, the model to be fitted is specified by a model formula involving the table variables. The MASS package provides a `coef()` method for "loglm" objects that extracts the estimated parameters and a `residuals()` method that calculates various types of residuals according to a `type` argument, one of "deviance", "pearson", "response". `vcd` and `vcdExtra` provide a variety of plotting methods, including `assoc()`, `sieve()`, `mosaic()` and `mosaic3d()` for "loglm" objects.

{ex:berkeley5}

EXAMPLE 8.1: Berkeley admissions

The UCBA admissions on admissions to the six largest graduate departments at U.C. Berkeley was examined using graphical methods in Chapter 4 (Example 4.14) and in Chapter 5 (Example 5.13). We can fit and compare several loglinear models as shown below.




Figure 8.1: Standard errors of residuals, $\sqrt{1 - h_i}$ decrease with expected frequencies. This plot shows why ordinary Pearson and deviance residuals may be misleading. The symbol size in the plot is proportional to leverage, h_i . Labels abbreviate Department, Gender and Admit, colored by Admit.

The model of mutual independence, $[A][D][G]$, is not substantively reasonable here, because the association of *Dept* and *Gender* should be taken into account to control for these variables, but we show it here to illustrate the form of the printed output, giving the Pearson χ^2 and likelihood-ratio G^2 tests of goodness of fit, as well as some optional arguments for saving additional components in the result.

```
data("UCBAdmissions")
library(MASS)
berk.loglm0 <- loglm(~ Dept + Gender + Admit, data=UCBAdmissions,
                    param=TRUE, fitted=TRUE)
berk.loglm0

## Call:
## loglm(formula = ~Dept + Gender + Admit, data = UCBAdmissions,
##       param = TRUE, fitted = TRUE)
##
## Statistics:
##               X^2 df P(> X^2)
## Likelihood Ratio 2098 16      0
## Pearson          2000 16      0
```

The argument `param=TRUE` stores the estimated parameters in the loglinear model and `fitted=TRUE` stores the fitted frequencies \hat{m}_{ijk} . The fitted frequencies can be extracted from the model object using `fitted()`.

```
structable(Dept ~ Admit+Gender, fitted(berk.loglm0))

##           Dept      A      B      C      D      E      F
```

```
## Admit      Gender
## Admitted Male      215.10 134.87 211.64 182.59 134.64 164.61
##           Female    146.68  91.97 144.32 124.51  91.81 112.25
## Rejected Male      339.63 212.95 334.17 288.30 212.59 259.91
##           Female    231.59 145.21 227.87 196.59 144.96 177.23
```

Similarly, you can extract the estimated parameters with `coef(berk.loglm0)`, and the Pearson residuals with `residuals(berk.loglm0, type="pearson")`.

Next, consider the model of conditional independence of gender and admission given department, $[AD][GD]$ that allows associations of admission with department and gender with department.

```
# conditional independence in UCB admissions data
berk.loglm1 <- loglm(~ Dept * (Gender + Admit), data=UCBAdmissions)
berk.loglm1

## Call:
## loglm(formula = ~Dept * (Gender + Admit), data = UCBAdmissions)
##
## Statistics:
##              X^2 df P(> X^2)
## Likelihood Ratio 21.74  6 0.001352
## Pearson          19.94  6 0.002840
```

Finally for this example, the model of homogeneous association, $[AD][AG][GD]$ can be fit as follows.²

```
berk.loglm2 <- loglm(~(Admit + Dept + Gender)^2, data=UCBAdmissions)
berk.loglm2

## Call:
## loglm(formula = ~(Admit + Dept + Gender)^2, data = UCBAdmissions)
##
## Statistics:
##              X^2 df P(> X^2)
## Likelihood Ratio 20.20  5 0.001144
## Pearson          18.82  5 0.002074
```

Neither of these models fits particularly well, as judged by the goodness-of-fit Pearson χ^2 and likelihood-ratio G^2 test against the saturated model. The `anova()` method for a nested collection of "loglm" models gives a series of likelihood-ratio tests of the difference, ΔG^2 between each sequential pair of models according to Eqn. (8.9).

```
anova(berk.loglm0, berk.loglm1, berk.loglm2, test="Chisq")

## LR tests for hierarchical log-linear models
##
## Model 1:
## ~Dept + Gender + Admit
## Model 2:
## ~Dept * (Gender + Admit)
## Model 3:
```

²It is useful to note here that the added term $[AG]$ allows a general association of admission with gender (controlling for department). A significance test for this term, or for model `berk.loglm2` against `berk.loglm1` is a proper test for the assertion of gender bias in admissions.

```
## ~(Admit + Dept + Gender)^2
##
##           Deviance df Delta(Dev) Delta(df) P(> Delta(Dev))
## Model 1      2097.67 16
## Model 2       21.74  6    2075.936      10      0.00000
## Model 3       20.20  5      1.531       1      0.21593
## Saturated      0.00  0     20.204       5      0.00114
```

The conclusion from these results is that the model `berk.loglm1` is not much worse than model `berk.loglm2`, but there is still significant lack-of-fit. The next example, using `glm()`, shows how to visualize the lack of fit and account for it.

△

8.3.5 Using `glm()`

{sec:loglin-glm}

Loglinear models fit with `glm()` require the data in a data frame in frequency form, for example as produced by `as.data.frame()` from a table. The model formula expresses the model for the frequency variable, using `family=poisson`.

{ex:berkeley6}

EXAMPLE 8.2: Berkeley admissions

For the $2 \times 2 \times 6$ `UCBAdmissions` table, first transform this to a frequency data frame:

```
berkeley <- as.data.frame(UCBAdmissions)
head(berkeley)

##      Admit Gender Dept Freq
## 1 Admitted   Male    A   512
## 2 Rejected   Male    A   313
## 3 Admitted Female    A    89
## 4 Rejected Female    A    19
## 5 Admitted   Male    B   353
## 6 Rejected   Male    B   207
```

Then, the model of conditional independence corresponding to `berk.loglm1` can be fit using `glm()` as shown below.

```
berk.glm1 <- glm(Freq ~ Dept * (Gender+Admit),
                 data=berkeley, family="poisson")
```

Similarly, the all two-way model of homogeneous association is fit using

```
berk.glm2 <- glm(Freq ~ (Dept + Gender + Admit)^2,
                 data=berkeley, family="poisson")
```

These models are equivalent to those fit using `loglm()` in Example 8.1. We get the same residual G^2 as before, and the likelihood-ratio test of ΔG^2 given by `anova()` gives the same result, that the model `berk.glm2` offers no significant improvement over model `berk.glm1`.

```
anova(berk.glm1, berk.glm2, test="Chisq")
```

```
## Analysis of Deviance Table
##
## Model 1: Freq ~ Dept * (Gender + Admit)
## Model 2: Freq ~ (Dept + Gender + Admit)^2
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1         6      21.7
## 2         5      20.2  1      1.53    0.22
```

Among other advantages of using `glm()` as opposed to `loglm()` is that an `anova()` method is available for *individual* "glm" models, giving significance tests of the contributions of each *term* in the model, as opposed to the tests for individual coefficients provided by `summary()`.³

```
anova(berk.glml, test="Chisq")

## Analysis of Deviance Table
##
## Model: poisson, link: log
##
## Response: Freq
##
## Terms added sequentially (first to last)
##
##
##              Df Deviance Resid. Df Resid. Dev Pr(>Chi)
## NULL                23      2650
## Dept                 5      160    18      2491 <2e-16 ***
## Gender                1      163    17      2328 <2e-16 ***
## Admit                 1      230    16      2098 <2e-16 ***
## Dept:Gender           5     1221    11       877 <2e-16 ***
## Dept:Admit            5      855     6        22 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

We proceed to consider what is wrong with these models and how they can be improved. A mosaic display can help diagnose the reason(s) for lack of fit of these models. We focus here on the model $[AD][GD]$ that allows an association between gender and department (i.e., men and women apply at different rates to departments).

The `mosaic()` method for "glm" objects in `vcdExtra` provides a `residuals_type` argument, allowing `residuals_type="rstandard"` for standardized residuals. The formula argument here pertains to the order of the variables in the mosaic, not a model formula.

```
library(vcdExtra)
mosaic(berk.glml, shade=TRUE, formula=~Admit+Dept+Gender,
       residuals_type="rstandard", labeling=labeling_residuals,
       main="Model: [AdmitDept][GenderDept]")
```

The mosaic display, shown in Figure 8.2, indicates that this model fits well (residuals are small) except in Department A. This suggests a model which allows an association between Admission and Gender in Department A only,

$$\log m_{ijk} = \mu + \lambda_i^A + \lambda_j^D + \lambda_k^G + \lambda_{ij}^{AD} + \lambda_{jk}^{DG} + I(j=1)\lambda_{ik}^{AG}, \quad (8.13) \quad \{\text{eq:berk2}\}$$

³Unfortunately, in the historical development of R, the `anova()` methods for linear and generalized linear models provide only *sequential* ("Type I") tests that are computationally easy, but useful only under special circumstances. The `car` package provides an analogous `Anova()` method that gives more generally useful *partial* ("Type II") tests for the additional contribution of each term beyond the others, taking marginal relations into account.

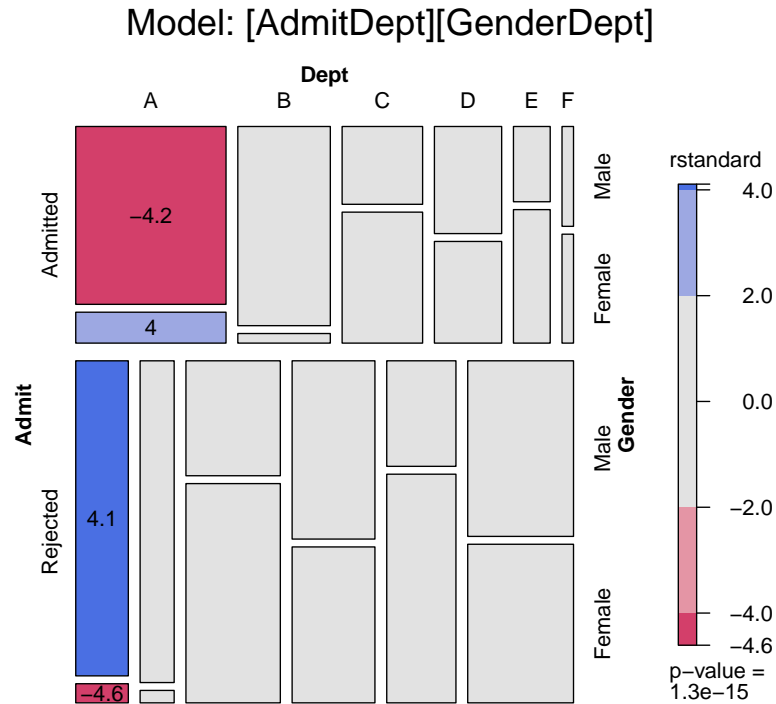


Figure 8.2: Mosaic display for the model [AD][GD], showing standardized residuals for the cell contributions to G^2 | Fig:berk-glm1-mosaic

where the indicator function $I(j = 1)$ equals 1 for Department A ($j = 1$) and is zero otherwise. This model asserts that Admission and Gender are conditionally independent, given Department, except in Department A. It has one more parameter than the conditional independence model, $[AD][GD]$, and forces perfect fit in the four cells for Department A.

Model Eqn. (8.13) may be fit with `glm()` by constructing a variable equal to the interaction of *gender* and *admit* with a dummy variable having the value 1 for Department A and 0 for other departments.

```
berkeley <- within(berkeley,
  dept1AG <- (Dept=='A') * (Gender=='Female') * (Admit=='Admitted'))
head(berkeley)

##      Admit Gender Dept Freq dept1AG
## 1 Admitted   Male   A   512        0
## 2 Rejected   Male   A   313        0
## 3 Admitted  Female   A    89        1
## 4 Rejected  Female   A    19        0
## 5 Admitted   Male   B   353        0
## 6 Rejected   Male   B   207        0
```

Fitting this model with the extra term `dept1AG` gives `berk.glm3`

```
berk.glm3 <- glm(Freq ~ Dept * (Gender+Admit) + dept1AG,
  data=berkeley, family="poisson")
```

This model does indeed fit well, and represents a substantial improvement over model `berk.glm1`:

```
vcdExtra::summarise(berk.glm3)

## Model Summary:
##           LR Chisq Df Pr(>Chisq)    AIC    BIC
## berk.glm3      2.68  5      0.749 -7.32 -39.4

anova(berk.glm1, berk.glm3, test="Chisq")

## Analysis of Deviance Table
##
## Model 1: Freq ~ Dept * (Gender + Admit)
## Model 2: Freq ~ Dept * (Gender + Admit) + dept1AG
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1         6      21.74
## 2         5       2.68  1      19.1  1.3e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The parameter estimate for the dept1AG term, $\hat{\lambda}_{ik}^{AG} = 1.052$ may be interpreted as the log odds ratio of admission for females as compared to males in Dept. A. The odds ratio is $\exp(1.052) = 2.86$, the same as the value calculated from the raw data (see Section 4.4.3).

```
coef(berk.glm3)[["dept1AG"]]

## [1] 1.052

exp(coef(berk.glm3)[["dept1AG"]])

## [1] 2.864
```

Finally, Figure 8.3 shows the mosaic for this revised model. The absence of shading indicates a well-fitting model.

```
mosaic(berk.glm3, shade=TRUE, formula=~Admit+Dept+Gender,
        residuals_type="rstandard", labeling=labeling_residuals,
        main="Model: [DeptGender][DeptAdmit] + DeptA*[GA]")
```

△

8.4 Equivalent logit models

{sec:loglin-logit}

Because loglinear models are formulated as models for the log (expected) frequency, they make no distinction between response and explanatory variables. In effect, they treat all variables as responses and describe their associations.

Logit (logistic regression) models, on the other hand, describe how the log odds for one variable depends on other, explanatory variables. There is a close connection between the two: When there is a response variable, each logit model for that response is equivalent to a loglinear model.

This relationship often provides a simpler way to formulate and test the model, and to plot and interpret the fitted results. Even when there is no response variable, the logit representation for one variable helps to interpret a loglinear model in terms of odds ratios. The price paid for this simplicity is that associations among the explanatory variables are not expressed in the model.

Model: [DeptGender][DeptAdmit] + DeptA*[GA]

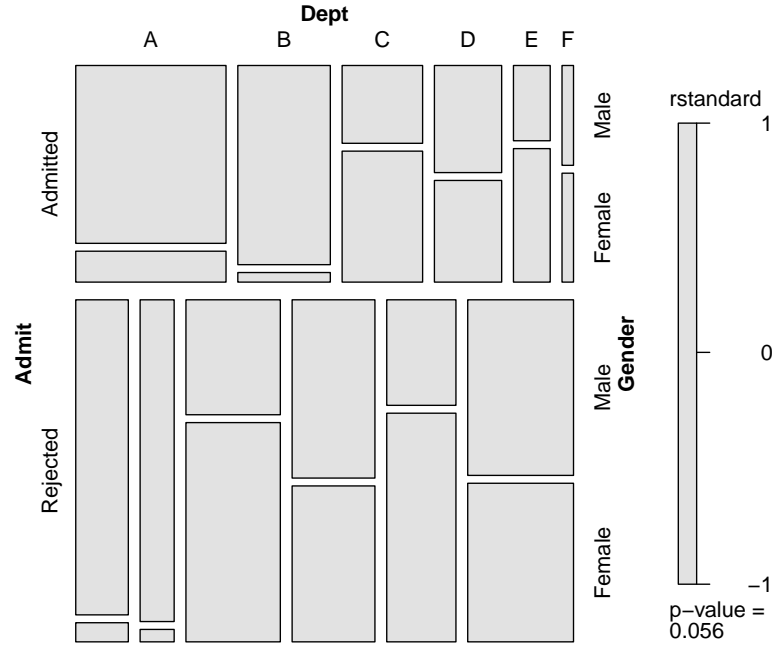


Figure 8.3: Mosaic display for the model `berk.glm3`, allowing an association of gender and admission in Department A.

Consider, for example, the model of homogeneous association, $[AB][AC][BC]$, Eqn. (8.5) for a three-way table, and let variable C be a binary response. Under this model, the logit for variable C is

$$\begin{aligned} L_{ij} = \log \left(\frac{\pi_{ij|1}}{\pi_{ij|2}} \right) &= \log \left(\frac{m_{ij1}}{m_{ij2}} \right) \\ &= \log(m_{ij1}) - \log(m_{ij2}) . \end{aligned}$$

Substituting from Eqn. (8.5), all terms which do not involve variable C cancel, and we are left with

$$\begin{aligned} L_{ij} = \log(m_{ij1}/m_{ij2}) &= (\lambda_1^C - \lambda_2^C) + (\lambda_{i1}^{AC} - \lambda_{i2}^{AC}) + (\lambda_{j1}^{BC} - \lambda_{j2}^{BC}) \\ &= 2\lambda_1^C + 2\lambda_{i1}^{AC} + 2\lambda_{j1}^{BC} , \end{aligned} \quad (8.14)$$

because all λ terms sum to zero. We are interested in how these logits depend on A and B , so we can simplify the notation by replacing the λ parameters with more familiar ones, $\alpha = 2\lambda_1^C$, $\beta_i^A = 2\lambda_{i1}^{AC}$, etc., which express this relation more directly,

$$L_{ij} = \alpha + \beta_i^A + \beta_j^B . \quad (8.15)$$

In the logit model Eqn. (8.15), the response, C , is affected by both A and B , which have additive effects on the log odds of response category C_1 compared to C_2 . The terms β_i^A and β_j^B correspond directly to $[AC]$ and $[BC]$ in the loglinear model Eqn. (8.5). The association among

the explanatory variables, $[AB]$ is assumed in the logit model, but this model provides no explicit representation of that association. The logit model Eqn. (8.14) is equivalent to the loglinear model $[AB][AC][BC]$ in goodness-of-fit and fitted values, and parameters in the two models correspond directly.

Table 8.1: Equivalent loglinear and logit models for a three-way table, with C as a binary response variable

Loglinear model	Logit model	Logit formula
$[AB][C]$	α	$C \sim 1$
$[AB][AC]$	$\alpha + \beta_i^A$	$C \sim A$
$[AB][BC]$	$\alpha + \beta_j^B$	$C \sim B$
$[AB][AC][BC]$	$\alpha + \beta_i^A + \beta_j^B$	$C \sim A + B$
$[ABC]$	$\alpha + \beta_i^A + \beta_j^B + \beta_{ij}^{AB}$	$C \sim A * B$

{tab:loglin-logit}

Table 8.1 shows the equivalent relationships between all loglinear and logit models for a three-way table when variable C is a binary response. Each model necessarily includes the $[AB]$ association involving the predictor variables. The most basic model, $[AB][C]$, is the intercept-only model, asserting constant odds for variable C . The saturated loglinear model $[ABC]$, allows an interaction in the effects of A and B on C , meaning that the AC association or odds ratio varies with B .

More generally, when there is a binary response variable, say R , and one or more explanatory variables, A, B, C, \dots , any logit model for R has an equivalent loglinear form. Every term in the logit model, such as β_{ik}^{AC} , corresponds to an association of those factors with R , that is, $[ACR]$ in the equivalent loglinear model.

The equivalent loglinear model must also include all associations among the explanatory factors, the term $[ABC\dots]$. Conversely, any loglinear model which includes all associations among the explanatory variables has an equivalent logit form. When the response factor has more than two categories, models for generalized logits (Section 7.5.4) also have an equivalent loglinear form.

{ex:berkeley7}

EXAMPLE 8.3: Berkeley admissions

The homogeneous association model, $[AD][AG][DG]$ did not fit the `UCBAdmissions` data very well, and we saw that the term $[AG]$ was unnecessary. Nevertheless, it is instructive to consider the equivalent logit model. We illustrate the features of the logit model which lead to the same conclusions and simplified interpretation from graphical displays.

Because Admission is a binary response variable, model Eqn. (8.6) is equivalent to the logit model,

$$L_{ij} = \log \left(\frac{m_{\text{Admit}(ij)}}{m_{\text{Reject}(ij)}} \right) = \alpha + \beta_i^{\text{Dept}} + \beta_j^{\text{Gender}} . \quad (8.16)$$

{eq:berk3}

That is, the logit model Eqn. (8.16) asserts that department and gender have additive effects on the log odds of admission. A significance test for the term β_j^{Gender} here is equivalent to the test of the $[AG]$ term for gender bias in the loglinear model. The observed log odds of admission here can be calculated as:

```
(obs <- log(UCBAdmissions[1,,] / UCBAdmissions[2,,]))
```

```
##           Dept
```

## Gender		A	B	C	D	E	F
## Male	0.4921	0.5337	-0.5355	-0.704	-0.957	-2.770	
## Female	1.5442	0.7538	-0.6604	-0.622	-1.157	-2.581	

With the data in the form of the frequency data frame `berkeley` we used in Example 8.2, the logit model Eqn. (8.16) can be fit using `glm()` as shown below. In the model formula, the binary response is `Admit=="Admitted"`. The `weights` argument gives the frequency, `Freq` in each table cell. **TODO: This gives the same fitted values, but not the same LR tests. Why?**

```
berk.logit2 <- glm(Admit=="Admitted" ~ Dept + Gender,
                  data=berkeley, weights=Freq, family="binomial")
summary(berk.logit2)
```

```
##
## Call:
## glm(formula = Admit == "Admitted" ~ Dept + Gender, family = "binomial",
##     data = berkeley, weights = Freq)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -25.342  -13.058   -0.163   16.017   21.320
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.5821    0.0690   8.44 <2e-16 ***
## DeptB         -0.0434    0.1098  -0.40  0.69
## DeptC         -1.2626    0.1066 -11.84 <2e-16 ***
## DeptD         -1.2946    0.1058 -12.23 <2e-16 ***
## DeptE         -1.7393    0.1261 -13.79 <2e-16 ***
## DeptF         -3.3065    0.1700 -19.45 <2e-16 ***
## GenderFemale    0.0999    0.0808   1.24  0.22
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6044.3  on 23  degrees of freedom
## Residual deviance: 5187.5  on 17  degrees of freedom
## AIC: 5201
##
## Number of Fisher Scoring iterations: 6
```

As in logistic regression models, parameter estimates may be interpreted as increments in the log odds, or $\exp(\beta)$ may be interpreted as the multiple of the odds associated with the explanatory categories. Because `glm()` uses a baseline category parameterization (by default) the coefficients of the first category of `Dept` and `Gender` are set to zero. You can see from the `summary()` output that the coefficients for the departments decline steadily from A–F.⁴ The coefficient $\beta_F^{\text{Gender}} = 0.0999$ for females indicates that, overall, women were $\exp(0.0999) = 1.105$ times as likely as male applicants to be admitted to graduate school at U.C. Berkeley, a 10% advantage.

Similarly, the logit model equivalent of the loglinear model Eqn. (8.13) `berk.glm3` containing the extra 1 df term for an effect of gender in Department A is

$$\{eq:berk4\} \quad L_{ij} = \alpha + \beta_i^{\text{Dept}} + I(j=1)\beta^{\text{Gender}}. \quad (8.17)$$

⁴In fact, the departments were labeled A–F in decreasing order of rate of admission.

This model can be fit as follows:

```
berkeley <- within(berkeley,
  dept1AG <- (Dept=='A')*(Gender=='Female'))
berk.logit3 <- glm(Admit=="Admitted" ~ Dept + Gender + dept1AG,
  data=berkeley, weights=Freq, family="binomial")
```

In contrast to the tests for individual coefficients, the `Anova()` method in the `car` package gives likelihood-ratio tests of the terms in a model. As mentioned earlier, this provides *partial* (“Type II”) tests for the additional contribution of each term beyond all others. **TODO: Fix: Too few digits are shown here for LR chisq**

```
library(car)
Anova(berk.logit2)

## Analysis of Deviance Table (Type II tests)
##
## Response: Admit == "Admitted"
##      LR Chisq Df Pr(>Chisq)
## Dept      763  5    <2e-16 ***
## Gender       2  1     0.22
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Anova(berk.logit3)

## Analysis of Deviance Table (Type II tests)
##
## Response: Admit == "Admitted"
##      LR Chisq Df Pr(>Chisq)
## Dept      647  5    < 2e-16 ***
## Gender       0  1     0.72
## dept1AG      18  1    2.7e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Plotting logit models

Logit models are easier to interpret than the corresponding loglinear models because there are fewer parameters, and because these parameters pertain to the odds of a response category rather than to cell frequency. Nevertheless, interpretation is often easier still from a graph than from the parameter values.

The simple interpretation of these logit models can be seen by plotting the logits for a given model. To do that, it is necessary to construct a data frame containing the observed (`obs`) and fitted (`fit`) for the combinations of gender and department.

```
pred2 <- cbind(berkeley[,1:3], fit=predict(berk.logit2))
pred2 <- cbind(subset(pred2, Admit=="Admitted"), obs=as.vector(obs))
head(pred2)

##      Admit Gender Dept      fit      obs
## 1 Admitted  Male   A   0.5821  0.4921
## 3 Admitted Female  A   0.6819  1.5442
## 5 Admitted  Male   B   0.5387  0.5337
## 7 Admitted Female  B   0.6385  0.7538
## 9 Admitted  Male   C  -0.6805 -0.5355
## 11 Admitted Female  C  -0.5807 -0.6604
```

In this form, these results can be plotted as a line plot of the fitted logits vs. department, with separate curves for males and females, and adding points to show the observed values. Here, we use `ggplot2` as shown below, with the `aes()` arguments `group=Gender`, `color=Gender`. This produces the left panel in Figure 8.4. The same steps for the model `berk.logit3` gives the right panel in this figure. The observed logits, of course, are the same in both plots.

```
library(ggplot2)
ggplot(pred2, aes(x=Dept, y=fit, group=Gender, color=Gender)) +
  geom_line(size=1.2) +
  geom_point(aes(x=Dept, y=obs, group=Gender, color=Gender), size=4) +
  ylab("Log odds (Admitted)") + theme_bw() +
  theme(legend.position=c(.8, .9),
        legend.title=element_text(size=14),
        legend.text=element_text(size=14))
```




Figure 8.4: Observed (points) and fitted (lines) log odds of admissions in the logit models for the UCBA admissions data. Left: the logit model Eqn. (8.16) corresponding to the loglinear model $[AD][AG][DG]$. Right: the logit model Eqn. (8.17), allowing only a 1 df term for Department A.

The effects seen in our earlier analyses (Examples 5.13, 5.14 and 8.2) may all be observed in these plots. In the left panel of Figure 8.4, corresponding to the loglinear model $[AD][AG][DG]$, the effect of gender, β_j^{Gender} , in the equivalent logit model is shown by the constant separation between the two curves. From the plot we see that this effect is very small (and nonsignificant). In the right panel, corresponding to the logit model Eqn. (8.17), there is no effect of gender on admission, except in department A, where the extra parameter allows perfect fit.

△

8.5 Zero frequencies

Cells with frequencies of zero create problems for loglinear and logit models. For loglinear models, most of the derivations of expected frequencies by maximum likelihood and other quantities

that depend on these (e.g., G^2 tests) assume that all $n_{ijk\dots} > 0$. In analogous logit models, the observed log odds (e.g., for a three-way table), $\log(n_{ij1}/n_{ij2})$, will be undefined if either frequency is zero.

Zero frequencies may occur in contingency tables for two different reasons:

- **structural zeros** (also called *fixed zeros*) will occur when it is impossible to observe values for some combinations of the variables. For these cases we should have $\hat{m}_i = 0$ whenever $n_i = 0$. For example, suppose we have three different methods of contacting people at risk for some obscure genetically inherited disease: newspaper advertisement, telephone campaign, and radio appeal. If each person contacted in any way is classified dichotomously by the three methods of contact, there can never be a non-zero frequency in the ‘No-No-No’ cell.⁵ Similarly, in a tabulation of seniors by gender and health concerns, there can never be males citing menopause or females citing prostate cancer. Square tables, such as wins and losses for sporting teams often have structural zeros in the main diagonal.
- **sampling zeros** (also called *random zeros*) occur when the total size of the sample is not large enough in relation to the probabilities in each of the cells to assure that someone will be observed in every cell. Here, it is permissible to have $\hat{m}_i > 0$ when $n_i = 0$. This problem increases with the number of table variables. For example, in a European survey of religious affiliation, gender and occupation, we may not happen to observe any female Muslim vineyard-workers in France, although such individuals surely exist in the population. Even when zero frequencies do not occur, tables with many cells relative to the total frequency tend to produce small expected frequencies in at least some cells, which tends to make the G^2 statistics for model fit and likelihood-ratio statistics for individual terms unreliable.

Following Birch (1963), Haberman (1974) and many others (e.g., Bishop *et al.*, 1975) identified conditions under which the maximum likelihood estimate for a given loglinear model does not exist, meaning that the algorithms used in `loglin()` and `glm()` do not converge to a solution. The problem depends on the number and locations of the zero cells, but not on the size of the frequencies in the remaining cells. Fienberg and Rinaldo (2007) give a historical overview of the problem and current approaches and Agresti (2013, §10.6) gives a compact summary.

In R, the mechanism to handle structural zeros in the IPF approach of `loglin()` and `loglm()` is to supply the argument `start`, giving a table conforming to the data, containing values of 0 in the locations of the zero cells, and non-zero elsewhere.⁶ In the `glm()` approach, the argument `subset=Freq > 0` can be used to remove the cells with zero frequencies from the data, or else, zero frequencies can be set to NA. This usually provides the correct degrees of freedom, however some estimated coefficients may be infinite.

For a complete table, the residual degrees of freedom are determined as

$$df = \# \text{ of cells} - \# \text{ of fitted parameters}$$

For tables with structural zeros, an analogous general formula is

$$df = (\# \text{ cells} - \# \text{ of parameters}) - (\# \text{ zero cells} - \# \text{ of NA parameters}) \quad (8.18) \quad \{\text{eq:dfzeros}\}$$

⁵Yet, if we fit an unsaturated model, expected frequencies may be estimated for all cells, and provide a means to estimate the total number at risk in the population. See Lindsey (1995, Section 5.4).

⁶If structural zeros are present, the calculation of degrees of freedom may not be correct. `loglm()` deducts one degree of freedom for each structural zero, but cannot make allowance for patterns of zeros based on the fitted margins that lead to gains in degrees of freedom due to smaller dimension in the parameter space. `loglin()` makes no such correction.

where NA parameters refers to parameters that cannot be estimated due to zero marginal totals in the model formula.

In contrast, sampling zeros are often handled by some modification of the data frequencies to ensure all non-zero cells. Some suggestions are:

- Add a small positive quantity (0.5 is often recommended) to *every* cell in the contingency table (Goodman, 1970), as is often done in calculating empirical log odds (Example 8.12); this simple approach over-smooths the data for unsaturated models, and should be deprecated, although widely used in practice.
- Replace sampling zeros by some small number, typically 10^{-10} or smaller (Agresti, 1990).
- Add a small quantity, like 0.1, to *all* zero cells, sampling or structural (Evers and Namboordiri., 1977).

In complex, sparse tables, a sensitivity analysis, comparing different approaches can help determine if the substantive conclusions vary with the approach to zero cells.

{ex:health}

EXAMPLE 8.4: Health concerns of teenagers

Fienberg (1980, Table 8-3) presented a classic example of structural zeros in the analysis of the $4 \times 2 \times 2$ table shown in Table 8.2. The data come from a survey of health concerns among teenagers, originally from Brunswick (1971). Among the health concerns, the two zero entries for menstrual problems among males are clearly structural zeros and there therefore one structural zero in the concern by gender marginal table. As usual, we abbreviate the table variables concern, age, gender by their initial letters, C, A, G below.

Table 8.2: Results from a survey of teenagers, regarding their health concerns. *Source:* Fienberg (1980, Table 8-3)

{tab:health}

Health Concerns	Gender:		Male		Female	
	Age:		12-15	16-17	12-15	16-17
sex, reproduction			4	2	9	7
menstrual problems			0	0	4	8
how healthy I am			42	7	19	10
nothing			57	20	71	21

The Health data is created as a frequency data frame as follows.

```
Health <- expand.grid(concerns = c("sex", "menstrual",
                                   "healthy", "nothing"),
                     age       = c("12-15", "16-17"),
                     gender    = c("M", "F"))
Health$Freq <- c(4, 0, 42, 57, 2, 0, 7, 20,
                 9, 4, 19, 71, 7, 8, 10, 21)
```

In this form, we first use `glm()` to fit two small models, neither of which involves the $\{CG\}$ margin. Model `health.glm0` is the model of mutual independence, $[C][A][G]$. Model `health.glm1` is the model of joint independence, $[C][AG]$, allowing an association between age and gender, but neither with concern. As noted above, the argument `subset=(Freq>0)` is used to eliminate the structural zero cells.

```
health.glm0 <-glm(Freq ~ concerns + age + gender, data=Health,
                  subset=(Freq>0), family=poisson)
health.glm1 <-glm(Freq ~ concerns + age * gender, data=Health,
                  subset=(Freq>0), family=poisson)
```

Neither of these fits the data well. To conserve space, we show only the results of the G^2 tests for model fit.

```
vcdExtra::summarise(glmlist(health.glm0, health.glm1))

## Model Summary:
##          LR Chisq Df Pr(>Chisq)  AIC   BIC
## health.glm0    27.7   8  0.000526 11.7 -17.4
## health.glm1    24.9   7  0.000795 10.9 -14.6
```

To see why, Figure 8.5 shows the mosaic display for model `health.glm1`, $[C][AG]$. Note that `mosaic()` takes care to make cells of zero frequency more visible by marking them with a small “o”, as these have an area of zero.

```
mosaic(health.glm1, ~concerns+age+gender, residuals_type="rstandard")
```




Figure 8.5: Mosaic display for the Health data, model `health.glm1` fig:health-mosaic

This suggests that there are important associations at least between concern and gender ($[CG]$) and between concern and age ($[CA]$). These are incorporated into the next model:

```
health.glm2 <-glm(Freq ~ concerns*gender + concerns*age, data=Health,
                  subset=(Freq>0), family=poisson)
vcdExtra::summarise(health.glm2)

## Model Summary:
##          LR Chisq Df Pr(>Chisq)  AIC   BIC
## health.glm2    4.66   3    0.198 -1.34 -12.2
```


The degrees of freedom are correct here. Eqn. (8.18), with 2 zero cells and 1 NA parameter due to the zero in the $\{CG\}$ margin gives $df = (16 - 12) - (2 - 1) = 3$. The loss of one estimable parameter can be seen in the output from `summary`.

```
summary(health.glm2)

##
## Call:
## glm(formula = Freq ~ concerns * gender + concerns * age, family = poisson,
##      data = Health, subset = (Freq > 0))
##
## Deviance Residuals:
##      1      3      4      5      7      8      9     10     11     12
##  0.236  0.585 -0.173 -0.300 -1.202  0.302 -0.149  0.000 -0.795  0.158
##     13     14     15     16
##  0.176  0.000  1.348 -0.282
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      1.266      0.445   2.84  0.0045 **
## concernsmenstrual -0.860      0.586  -1.47  0.1425
## concernshealthy   2.380      0.471   5.05  4.4e-07 ***
## concernsnothing   2.800      0.462   6.07  1.3e-09 ***
## genderF           0.981      0.479   2.05  0.0405 *
## age16-17          -0.368      0.434  -0.85  0.3964
## concernsmenstrual:genderF      NA         NA      NA      NA
## concernshealthy:genderF -1.505      0.533  -2.82  0.0047 **
## concernsnothing:genderF -0.803      0.503  -1.60  0.1105
## concernsmenstrual:age16-17  1.061      0.750   1.41  0.1574
## concernshealthy:age16-17  -0.910      0.513  -1.77  0.0761 .
## concernsnothing:age16-17  -0.771      0.469  -1.64  0.1005
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 252.4670  on 13  degrees of freedom
## Residual deviance:  4.6611  on  3  degrees of freedom
## AIC: 87.66
##
## Number of Fisher Scoring iterations: 4
```

In contrast, `loglm()` reports the degrees of freedom incorrectly for models containing zeros in any fitted margin. For use with `loglm()`, we convert it to a $4 \times 2 \times$ table.

```
health.tab <- xtabs(Freq ~ concerns + age + gender, data = Health)
```

The same three models are fitted with `loglm()` as shown below. The locations of the positive frequencies are marked in the array `nonzeros` and supplied as the value of the `start` argument.

```
nonzeros <- ifelse(health.tab>0, 1, 0)
health.loglm0 <- loglm(~ concerns + age + gender,
  data = health.tab, start = nonzeros)
health.loglm1 <- loglm(~ concerns + age * gender,
  data = health.tab, start = nonzeros)
# df is wrong
health.loglm2 <- loglm(~ concerns*gender + concerns*age,
  data = health.tab, start = nonzeros)
vcdExtra::summarise(loglm1list(health.loglm0, health.loglm1, health.loglm2))

## Re-fitting to get fitted values
## Re-fitting to get fitted values
```

```
## Re-fitting to get fitted values
## Model Summary:
##      LR Chisq Df Pr(>Chisq)    AIC    BIC
## health.loglm0  27.74  8    0.0005  11.74 -17.36
## health.loglm1  24.89  7    0.0008  10.89 -14.58
## health.loglm2   4.66  2    0.0972   0.66  -6.62
```

The results agree with those of `glm()`, except for the degrees of freedom for the last model.

△

8.6 Models for ordinal variables

{sec:loglin-ordinal}

Standard loglinear models treat all classification variables as nominal, unordered factors. In these models, all statistical tests are identical and parameter estimates are equivalent if the categories of any of the table variable are reordered. Yet we have seen that the ordering of categories often provides important information about the nature of associations and we showed (Section 4.2.4) that non-parametric tests which take into account the ordered nature of a factor are more powerful.

Correspondence analysis plots (Chapter 6) make it easy to see the relationships between ordinal variables, because the method assigns quantitative scores to the table variables which maximally account for their association. As we saw for the hair-eye color data (Figure 6.1) and the mental impairment data (Figure 6.2), an association can be interpreted in terms of ordered categories when the points for two factors are ordered similarly, usually along the first CA dimension.

Similarly, in a mosaic display, an ordered associative effect is seen when the residuals have an opposite-corner pattern of positive and negative signs and magnitudes (e.g., for the hair-eye color data, Figure 5.4 or the Titanic data, Figure ??). In these cases loglinear and logit models which use the ordered nature of the factors offer several advantages.

- Because they are more focused, tests which use the ordinal structure of the table variables are more powerful when the association varies systematically with the ordered values of a factor.
- Because they consume fewer degrees of freedom, we can fit unsaturated models where the corresponding model for nominal factors would be saturated. In a two-way table, for example, a variety of models for ordinal factors may be proposed which are intermediate between the independence model and the saturated model.
- Parameter estimates from these models are fewer in number, are easier to interpret, and quantify the nature of effects better than corresponding quantities in models for nominal factors. Estimating fewer parameters typically gives smaller standard errors, as we saw in Example ??.

These advantages are analogous to the use of tests for trends or polynomial contrasts in ANOVA models. More importantly, in some research areas in the social sciences (where categorical data is commonplace), models for ordinal variables have proved crucial in theory construction and debates, giving more precise tests of hypotheses than available from less focused or descriptive methods (Agresti, 1984).

8.6.1 Loglinear models for ordinal variables

{sec:loglin}

For a two-way table, when either the row variable or the column variable, or both, are ordinal, one simplification comes from assigning ordered scores, $\mathbf{a} = \{a_i\}, a_1 \leq a_2 \leq \dots \leq a_I$, and/or $\mathbf{b} = \{b_j\}, b_1 \leq b_2 \leq \dots \leq b_J$ to the categories so that the ordinal relations are necessarily included in the model. Typically, equally spaced scores are used, for example, integer scores, $\{a_i\} = i$, or the zero-sum equivalent, $\{a_i\} = i - (I + 1)/2$ (e.g., $\{a_i\} = \{-1, 0, 1\}$ for $I = 3$).

Using such scores gives simple interpretations of the association parameters in terms of *local odds ratios* for adjacent 2×2 subtables,

{eq:loddsratios}

$$\theta_{ij} = \frac{m_{ij} m_{i+1,j+1}}{m_{i,j+1} m_{i+1,j}} , \quad (8.19)$$

which is the odds ratio for pairs of adjacent rows and adjacent columns.

When both variables are assigned scores, this gives the **linear-by-linear model** ($L \times L$)

{eq:linlin}

$$\log(m_{ij}) = \mu + \lambda_i^A + \lambda_j^B + \gamma a_i b_j . \quad (8.20)$$

Because the scores \mathbf{a} and \mathbf{b} are fixed, this model has only one extra parameter, γ , compared to the independence model, which is the special case, $\gamma = 0$. In contrast, the saturated model, allowing general association λ_{ij}^{AB} uses $(I - 1)(J - 1)$ additional parameters.

The terms $\gamma a_i b_j$ in Eqn. (8.20) describe a pattern of association where deviations from independence increase linearly with a_i and b_j in opposite directions towards the opposite corners of the table, as we have often observed in mosaic displays.

In the linear-by-linear association model, the local log odds ratios are

$$\log \theta_{ij} = \gamma(a_{i+1} - a_i)(b_{j+1} - b_j) ,$$

which reduces to

$$\log \theta_{ij} = \gamma$$

for integer-spaced scores, so γ is the common local log odds ratio. As a result, the linear-by-linear model is sometimes called the **uniform association model** (Goodman, 1979).

Generalizations of the linear-by-linear model result when only one variable is assigned scores. In the **row effects model** (R), the row variable, A , is treated as nominal, while the column variable, B , is assigned ordered scores $\{b_j\}$. The loglinear model is then

{eq:roweff}

$$\log(m_{ij}) = \mu + \lambda_i^A + \lambda_j^B + \alpha_i b_j , \quad (8.21)$$

where the α_i parameters are the *row effects*. An additional constraint, $\sum_i \alpha_i = 0$ or $\alpha_1 = 0$ is imposed, so that model Eqn. (8.21) has only $(I - 1)$ more parameters than the independence model. The linear-by-linear model is the special case where the row effects are equally spaced, and the independence model is the special case where all $\alpha_i = 0$.

The row-effects model Eqn. (8.21) also has a simple odds ratio interpretation. The local log odds ratio for adjacent pairs of rows and columns is

$$\log \theta_{ij} = \alpha_{i+1} - \alpha_i ,$$

which is constant for all pairs of adjacent columns. Plots of the local log odds ratio against i would appear as a set of parallel curves.

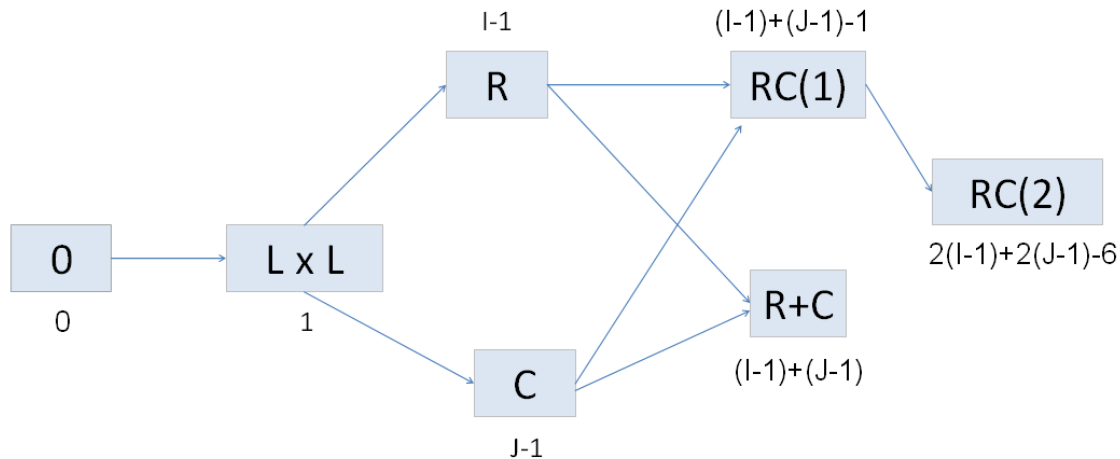


Figure 8.6: Nesting relationships among some association models for an $I \times J$ table specifying the association parameters, λ_{ij}^{AB} . Model **0** is the independence model. Formulas near the boxes give the number of identifiable association parameters. Arrows point from one nested model to another that is a more general version.

{fig:assoc-models}

In the analogous *column effects model* (C), $(J - 1)$ linearly independent column effect parameters β_j are estimated for the column variable, while fixed scores $\{a_i\}$ are assigned to the row variable. It is also possible to fit a *row plus column effects model* (R+C), that assigns specified scores to both the rows and column variables.

Nesting relationships among these models and others described in Section 8.6.2 are shown in Figure 8.6. Any set of models connected by a path can be directly compared with likelihood-ratio tests of the form $G^2(M_2|M_1)$.

In R, the $L \times L$, row effects and column effects models can all be fit using `glm()` simply by replacing the appropriate table factor variable(s) with their `as.numeric()` equivalents.

{ex:mental4}

EXAMPLE 8.5: Mental impairment and parents' SES

The `Mental` data on the mental health status of young New York residents in relation to their parents' socioeconomic status was examined in Example 4.6 using CMH tests for ordinal association and in Example 6.2 using correspondence analysis. Figure 6.2 showed that nearly all of the association in the table was accounted for by a single dimension along which both factors were ordered, consistent with the view that mental health increased in relation to parents' SES.

Because these models provide their interpretations in terms of local odds ratios, Eqn. (8.19), it is helpful to see these values for the observed data, corresponding to the saturated model. The values $\log(\theta_{ij})$ are calculated by `loddsratio()` in `vcdExtra`, with the data in table form.

```
(Mental.tab <- xtabs(Freq ~ mental+ses, data=Mental))
```

```
##          ses
## mental    1  2  3  4  5  6
## Well      64 57 57 72 36 21
## Mild      94 94 105 141 97 71
## Moderate  58 54 65 77 54 54
## Impaired  46 40 60 94 78 71
```

```
loddsratio(Mental.tab)
```

```
## log odds ratios for mental and ses
##
##
##          ses
## mental      1:2      2:3      3:4      4:5      5:6
## Well:Mild    0.11583 0.11067 0.06118 0.31910 0.22697
## Mild:Moderate -0.07146 0.07474 -0.12538 0.01923 0.31203
## Moderate:Impaired -0.06830 0.22006 0.27953 0.16824 -0.09403
```




Figure 8.7: Shaded-square plot of the local odds ratios in the `Mental` data.

A simple plot of these values, using area- and color-proportional shaded squares is shown in Figure 8.7. This plot is drawn using the `corrplot` package. It is easy to see that most of the local odds ratios are mildly positive.

```
M <- as.matrix(loddsratio(Mental.tab))
library(corrplot)
corrplot(M, method="square", is.corr=FALSE,
          tl.col="black", tl.srt=0, tl.offset=1)
```

For comparison with the $L \times L$ model fitted below, the mean local log odds ratio is 0.103.

```
mean(loddsratio(Mental.tab)$coefficients)

## [1] 0.1032
```

As a baseline, we first fit the independence model (testing $H_0 : \log(\theta_{ij}) = 0$) with `glm()`. As expected, this model fits quite badly, with $G^2(15) = 47.418$.

```
indep <- glm(Freq ~ mental + ses,
              family = poisson, data = Mental)
vcdExtra::summarise(indep)

## Model Summary:
##      LR Chisq Df Pr(>Chisq)  AIC   BIC
## indep    47.4 15  3.16e-05 17.4 -63.8
```

The mosaic display of standardized residuals from this model is shown in Figure 8.8. The argument `labeling=labeling_residuals` is used to show the numerical values in the cells with absolute values greater than `suppress=1`.

```
long.labels <- list(set_varnames = c(mental="Mental Health Status",
                                     ses="Parent SES"))
mosaic(indep,
        gp=shading_Friendly,
        residuals_type="rstandard",
        labeling_args = long.labels,
        labeling=labeling_residuals, suppress=1,
        main="Mental health data: Independence")
```




Figure 8.8: Mosaic display of the independence model for the mental health data. fig:mental-indep

This figure shows the classic opposite-corner pattern of the signs and magnitudes of the residuals that would arise if the association between mental health and SES was could be explained by the ordinal relation of these factors using one of the $L \times L$, R or C models.

To fit such ordinal models, you can use `as.numeric()` on a factor variable to assign integer scores, or assign other values if integer spacing is not appropriate.

```
Cscore <- as.numeric(Mental$ses)
Rscore <- as.numeric(Mental$mental)
```

Then, the $L \times L$, R and C models can be fit as follows, where beyond the main effects of mental and ses, their association is represented as the interaction of the numeric score(s) or factor(s), as appropriate in each case.

```
linlin <- glm(Freq ~ mental + ses + Rscore:Cscore,
              family = poisson, data = Mental)
roweff <- glm(Freq ~ mental + ses + mental:Cscore,
              family = poisson, data = Mental)
coleff <- glm(Freq ~ mental + ses + Rscore:ses,
              family = poisson, data = Mental)
```

Goodness-of-fit tests for these models are shown below. They show that all of the $L \times L$, R and C models are acceptable in terms of the likelihood-ratio G^2 . The $L \times L$ model, with only one more parameter than the independence model is judged the best by both AIC and BIC.

```
vcdExtra::summarise(glmlist(indep, linlin, roweff, coleff))

## Model Summary:
##      LR Chisq Df Pr(>Chisq)    AIC    BIC
## indep      47.4 15      0.000    17.4 -63.8
## linlin      9.9 14      0.770   -18.1 -93.9
## roweff      6.3 12      0.901   -17.7 -82.7
## coleff      6.8 10      0.741   -13.2 -67.3
```

In cases where such overall tests are unclear, you can carry out tests of nested sets of models using `anova()`, giving tests of ΔG^2 .

```
anova(indep, linlin, roweff, test="Chisq")

## Analysis of Deviance Table
##
## Model 1: Freq ~ mental + ses
## Model 2: Freq ~ mental + ses + Rscore:Cscore
## Model 3: Freq ~ mental + ses + mental:Cscore
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1         15         47.4
## 2          14          9.9  1      37.5    9e-10 ***
## 3           12          6.3  2       3.6     0.16
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

anova(indep, linlin, coleff, test="Chisq")

## Analysis of Deviance Table
##
## Model 1: Freq ~ mental + ses
## Model 2: Freq ~ mental + ses + Rscore:Cscore
## Model 3: Freq ~ mental + ses + Rscore:s ses
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1         15         47.4
## 2          14          9.9  1      37.5    9e-10 ***
## 3           10          6.8  4       3.1     0.55
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Under the $L \times L$ model, the estimate of the coefficient of `Rscore:Cscore` is $\hat{\gamma} = 0.0907$ (s.e.=0.015) with unit-spaced scores, as shown below.

```
# interpret linlin association parameter
coef(linlin)[["Rscore:Cscore"]]

## [1] 0.09069

exp(coef(linlin)[["Rscore:Cscore"]])

## [1] 1.095
```

This corresponds to a local odds ratio, $\hat{\theta}_{ij} = \exp(0.0907) = 1.095$. This single number describes the association succinctly: each step down the socioeconomic scale increases the odds of being classified one step poorer in mental health by 9.5%.

△

8.6.2 Log-multiplicative (RC) models

{sec:RCmodels}

The association models described above are all more parsimonious and easier to interpret than the saturated model. However, they depend on assigning fixed and possibly arbitrary scores to the variable categories. A generalization of the $L \times L$ model that treats *both* row and column scores as parameters is the **row-and-column effects model** (RC(1)) suggested by Goodman (1979),

$$\log(m_{ij}) = \mu + \lambda_i^A + \lambda_j^B + \gamma \alpha_i \beta_j, \quad (8.22) \quad \{\text{eq:RC1}\}$$

where γ , α and β comprise additional parameters to be estimated beyond the independence model.⁷ This model has a close connection with correspondence analysis (Goodman, 1985), where the estimated scores α and β are analogous to correspondence analysis scores on a first dimension.⁸

For identifiability and interpretation it is necessary to impose some normalization constraints on the α and β . An *unweighted, unit standardized* solution forces $\sum_i \alpha_i = \sum_j \beta_j = 0$ and $\sum_i \alpha_i^2 = \sum_j \beta_j^2 = 1$. Alternatively, and more akin to correspondence analysis solutions, the *marginally weighted* solution uses the marginal probabilities π_{i+} of the row variable and π_{+j} of the columns as weights.

$$\begin{aligned} \sum_i \alpha_i \pi_{i+} &= \sum_j \beta_j \pi_{+j} = 0 \\ \sum_i \alpha_i^2 \pi_{i+} &= \sum_j \beta_j^2 \pi_{+j} = 1 \end{aligned} \quad (8.23) \quad \{\text{eq:RC-constraints}\}$$

Goodman (1986) generalized this to multiple bilinear terms of the form $\gamma_k \alpha_{ik} \beta_{jk}$, and showed that all associations in the saturated model could be expressed as

$$\lambda_{ij}^{AB} = \sum_{k=1}^M \gamma_k \alpha_{ik} \beta_{jk} \quad M = \min(I-1, J-1). \quad (8.24) \quad \{\text{eq:RCm}\}$$

For example, an RC(2) model with two multiplicative terms is analogous to a two-dimensional correspondence analysis solution.

The RC model is *not* a loglinear model because it contains a multiplicative term in the parameters. This model and a wide variety of other nonlinear models for categorical data can be fit using `gnm()` in the `gnm` package.

{ex:mental5}

EXAMPLE 8.6: Mental impairment and parents' SES

The `gnm` package provides a number of functions that can be used in model formulas for nonlinear association terms. Among these, `Mult()` expresses a multiplicative association in terms

⁷In contrast to the R, C and R+C models, RC models do not assume that the categories are appropriately ordered because the category scores are estimated from the data.

⁸However, when estimated by maximum likelihood, the RC(1) model allows likelihood-ratio tests of parameters and model fit, while such significance tests are not available for correspondence analysis.

of two (or more) factors. The RC(1) model for factors A , B uses $\text{Mult}(A, B)$ for the association term in Eqn. (8.22). Multiple multiplicative RC terms, as in Eqn. (8.24) can be expressed using instances $(\text{Mult}(A, B), m)$.

To illustrate, we fit the RC(1) and RC(2) models to the `Mental` data using `gnm()`. In this table, both factors are ordered, but we don't want to use the default polynomial contrasts, so we set their contrast attributes to `treatment`.

```
library(gnm)
Mental$mental <- C(Mental$mental, treatment)
Mental$ses <- C(Mental$ses, treatment)
RC1 <- gnm(Freq ~ mental + ses + Mult(mental, ses),
           family = poisson, data = Mental, verbose=FALSE)
RC2 <- gnm(Freq ~ mental + ses + instances(Mult(mental, ses), 2),
           family = poisson, data = Mental, verbose=FALSE)
```

For comparison with the loglinear association models fit in Example 8.5 we show the G^2 goodness of fit tests for all these models. The ordinal loglinear models and the RC models all fit well, with the $L \times L$ model preferred on the basis of parsimony by AIC and BIC.

```
vcdExtra::summarise(glmlist(indep, linlin, roweff, coleff, RC1, RC2))
```

## Model Summary:	##	LR	Chisq	Df	Pr(>Chisq)	AIC	BIC
## indep		47.4	15		0.000	17.42	-63.8
## linlin		9.9	14		0.770	-18.10	-93.9
## roweff		6.3	12		0.901	-17.72	-82.7
## coleff		6.8	10		0.741	-13.17	-67.3
## RC1		3.6	8		0.894	-12.43	-55.7
## RC2		0.5	3		0.914	-5.48	-21.7

The substantive difference between the $L \times L$ model and the RC(1) model is whether the categories of mental health status and SES can be interpreted as equally spaced along some latent continua, versus the alternative that category spacing is unequal. We can test this directly using the likelihood-ratio test, $G^2(L \times L | RC(1))$. Similarly, model RC1 is nested within model RC2, so $G^2(RC(1) | RC(2))$ gives a direct test of the need for a second dimension.

```
anova(linlin, RC1, RC2, test="Chisq")
```

## Analysis of Deviance Table	##	## Model 1: Freq ~ mental + ses + Rscore:Cscore	## Model 2: Freq ~ mental + ses + Mult(mental, ses)	## Model 3: Freq ~ mental + ses + Mult(mental, ses, inst = 1) + Mult(mental, ses, inst = 2)	##	## Resid. Df	## Resid. Dev	## Df	## Deviance	## Pr(>Chi)
## 1		14	9.90							
## 2		8	3.57	6	6.32	0.39				
## 3		3	0.52	5	3.05	0.69				

We see that estimated scores for the categories in the model RC1 do not provide a significantly better fit, and there is even less evidence for a second dimension of category parameters in the RC2 model.

Nevertheless, for cases where RC models *do* provide some advantage, it is useful to know how to visualize the estimated category parameters. The key to this is the function `getContrasts()` which computes contrasts or scaled contrasts for a set of (non-eliminated) parameters from a

"gnm" model, together with standard errors for the estimated contrasts following the methods of Firth (2003), Firth and Menezes (2004). The details are explained in `help(getContrasts)` and in vignette("gnmOverview") that comes with the `gnm` package.

The coefficients in the marginally-weighted solution Eqn. (8.23) can be obtained as follows.

```
rowProbs <- with(Mental, tapply(Freq, mental, sum) / sum(Freq))
colProbs <- with(Mental, tapply(Freq, ses, sum) / sum(Freq))
mu <- getContrasts(RC1, pickCoef(RC1, "[.]mental"),
                  ref = rowProbs, scaleWeights = rowProbs)
nu <- getContrasts(RC1, pickCoef(RC1, "[.]ses"),
                  ref = colProbs, scaleWeights = colProbs)
```

In our notation, the coefficients α and β can be extracted as the `qvframe` component of the "qv" object returned by `getContrasts()`.

```
(alpha <- mu$qvframe)

##                               Estimate Std. Error
## Mult(., ses).mentalWell      -1.6738      0.1904
## Mult(., ses).mentalMild       -0.1401      0.2002
## Mult(., ses).mentalModerate    0.1367      0.2795
## Mult(., ses).mentalImpaired    1.4105      0.1742

(beta <- nu$qvframe)

##                               Estimate Std. Error
## Mult(mental, .).ses1 -1.11136      0.2992
## Mult(mental, .).ses2 -1.12046      0.3142
## Mult(mental, .).ses3 -0.37075      0.3192
## Mult(mental, .).ses4  0.02701      0.2733
## Mult(mental, .).ses5  1.00948      0.3147
## Mult(mental, .).ses6  1.81665      0.2810
```

For plotting this RC(1) solution for the scaled category scores together with their estimated standard errors, a `dotchart()`, shown in Figure 8.9 provides a reasonable visualization.

To create this plot, first combine the row and column scores in a data frame, and add columns `lower`, `upper` corresponding to ± 1 standard error (or some other multiple).

```
scores <- rbind(alpha, beta)
scores <- cbind(scores,
                factor=c(rep("mental", 4), rep("ses", 6)) )
rownames(scores) <- c(levels(Mental$mental), levels(Mental$ses))
scores$lower <- scores[,1]-scores[,2]
scores$upper <- scores[,1]+scores[,2]
scores

##           Estimate Std. Error factor  lower  upper
## Well      -1.67378      0.1904 mental -1.8642 -1.48336
## Mild      -0.14009      0.2002 mental -0.3403  0.06009
## Moderate   0.13669      0.2795 mental -0.1428  0.41617
## Impaired   1.41055      0.1742 mental  1.2364  1.58473
## 1          -1.11136      0.2992   ses -1.4106 -0.81215
## 2          -1.12046      0.3142   ses -1.4347 -0.80624
## 3          -0.37075      0.3192   ses -0.6899 -0.05160
## 4           0.02701      0.2733   ses -0.2463  0.30028
## 5           1.00948      0.3147   ses  0.6948  1.32418
## 6           1.81665      0.2810   ses  1.5357  2.09760
```




Figure 8.9: Dotchart of the scaled category scores for the RC(1) model fit the the mental health data. Error bars show ± 1 standard error.

{fig:mental}

The dotchart shown in Figure 8.9 is then a plot of Estimate, grouped by factor, with arrows showing the range of lower to upper for each parameter.

```
with(scores, {
  dotchart(Estimate, groups=factor, labels=rownames(scores),
    cex=1.2, pch=16, xlab="RC1 Score",
    xlim=c(min(lower), max(upper)))
  arrows(lower, c(8+(1:4), 1:6), upper, c(8+(1:4), 1:6),
    col="red", angle=90, length=.05, code=3, lwd=2)
})
```

In this plot, the main substantive difference from the $L \times L$ model is in the spacing of the lowest two categories of `ses` and the middle two categories of `mental` which are not seen to differ in the RC1 model.

The coefficients in the RC2 model can also be plotted (in a 2D plot) by extracting the coefficients from the "gnm" object and reshaping them to 2-column matrices. The function `pickCoef()` is handy here to get the indices of a subset of parameters by matching a pattern in their names.

```
alpha <- coef(RC2)[pickCoef(RC2, "[.]mental")]
alpha <- matrix(alpha, ncol=2)
rownames(alpha) <- levels(Mental$mental)
colnames(alpha) <- c("Dim1", "Dim2")
alpha

##          Dim1      Dim2
## Well      1.725 -0.40958
## Mild      0.217 -0.13369
```

```
## Moderate 0.122 0.32796
## Impaired -1.388 -0.09192

beta <- coef(RC2)[pickCoef(RC2, "[.]ses")]
beta <- matrix(beta, ncol=2)
rownames(beta) <- levels(Mental$ses)
colnames(beta) <- c("Dim1", "Dim2")
beta

##          Dim1      Dim2
## 1  0.18998  0.07009
## 2  0.19897  0.12271
## 3  0.07020  0.11509
## 4 -0.03064 -0.18906
## 5 -0.17627 -0.07526
## 6 -0.26913  0.48921
```

The simple, unweighted scaling to mean 0, variance 1 can be obtained with `scale()`:

```
alpha <- scale(alpha)
beta <- scale(beta)
```

Alternatively, the marginal-weighted scaling of Eqn. (8.23) is obtained by centering at the weighted mean and dividing by the weighted sum of squares. We use this scaling here.

```
alpha <- apply(alpha, 2, function(x) x - sum(x*rowProbs))
alpha <- apply(alpha, 2, function(x) x/sqrt(sum(x^2 * rowProbs)))
beta <- apply(beta, 2, function(x) x - sum(x*colProbs))
beta <- apply(beta, 2, function(x) x/sqrt(sum(x^2 * colProbs)))
```

To plot these category scores, first combine them into a single data frame,

```
scores <- data.frame(rbind(alpha,beta))
scores$factor <- c(rep("mental", 4), rep("ses", 6))
scores$probs <- c(rowProbs, colProbs)
scores

##          Dim1      Dim2 factor probs
## Well      1.61612 -1.39928 mental 0.1849
## Mild      0.11716 -0.24803 mental 0.3627
## Moderate  0.02277  1.67839 mental 0.2181
## Impaired -1.47795 -0.07374 mental 0.2343
## 1         1.14884  0.06220      ses 0.1578
## 2         1.20354  0.31752      ses 0.1476
## 3         0.42039  0.28056      ses 0.1729
## 4        -0.19294 -1.19538      ses 0.2313
## 5        -1.07857 -0.64313      ses 0.1596
## 6        -1.64335  2.09605      ses 0.1307
```

Then, we use `xyplot()` to plot the scores on Dim2 against Dim1, with separate lines and colors for the two factors. The resulting plot is shown in Figure 8.10.

```
library(lattice)
xyplot(Dim2 ~ Dim1, groups=factor, data=scores, type="b",
       cex=1.3, pch=16, lwd=2, aspect="iso",
       panel=function(x, y, ...) {
         panel.xyplot(x, y, ...)
       })
```




Figure 8.10: Scaled category scores for the RC(2) model fit to the mental health data.

{fig:mental}

```

panel.text(x=x, y=y, labels=rownames(scores), pos=1, cex=1.2)
panel.abline(h=0, col="gray")
panel.abline(v=0, col="gray")
}

```

)

The patterns of the row and column category scores here are quite similar to the 2D correspondence analysis solution shown in Figure 6.2. The main difference is in the relative scaling of the axes. In Figure 8.10, the variances of the two dimensions are equated; in the correspondence analysis plot, the axes are scaled in relation to their contributions to Pearson χ^2 , allowing an interpretation of distance between points in terms of χ^2 -distance.

△

8.7 Square tables

{sec:loglin-square}

Square tables, where the row and column variables have the same categories comprise an important special case for loglinear models that can account for associations more parsimoniously than the saturated model. Some examples are the data on visual acuity in Example 4.13, categorical ratings of therapy clients by two observers, and mobility tables, tracking the occupational categories between generations in the same families or migration tables, giving movement of people between regions. The latter topics have been important in sociological and geographic research and has spurred the development of a wide range of specialized loglinear models for this purpose.

8.7.1 Quasi-independence, symmetry, quasi-symmetry and topological models

c:sq-quasi}

In many square tables, such as the `Vision` data, independence is not a credible hypothesis because the diagonal cells, representing equal values of the row and column variables tend to be very large and often contribute most of the lack of fit. A substantively more interesting hypothesis is whether the table exhibits independence, ignoring the diagonal cells. This leads to what is called the **quasi-independence model**, that specifies independence only in the off-diagonal cells.

For a two-way table, quasi-independence can be expressed as

$$\pi_{ij} = \pi_{i+}\pi_{+j} \quad \text{for } i \neq j$$

or in loglinear form as

$$\log m_{ij} = \mu + \lambda_i^A + \lambda_j^B + \delta_i I(i = j) .$$

This model effectively adds one parameter, δ_i , for each main diagonal cell which fits those frequencies perfectly.

Another hypothesis of substantive interest for square tables, particularly those concerning occupational and geographical mobility is that the joint distribution of row and column variables is symmetric, that is, $\pi_{ij} = \pi_{ji}$ for all $i \neq j$. For example, this **symmetry model** (S) asserts that sons are as likely to move from their father's occupation i to another, j , as the reverse. This form of symmetry is quite strong, because it also implies **marginal homogeneity** (MH), that the marginal probabilities of the row and column variables are equal, $\pi_{i+} = \sum_j \pi_{ij} = \sum_j \pi_{ji} = \pi_{+i}$ for all i .

To separate marginal homogeneity from symmetry of the association terms per se, the model of **quasi-symmetry** (QS) uses the standard main-effect terms in the loglinear model,

$$\log m_{ij} = \mu + \lambda_i^A + \lambda_j^B + \lambda_{ij} , \quad (8.25) \quad \{\text{eq:quasi-symm}\}$$

where $\lambda_{ij} = \lambda_{ji}$. It can be shown (Causinus, 1966) that

$$\begin{aligned} \text{symmetry} &= \text{quasi-symmetry} + \text{marginal homogeneity} \\ G^2(S) &= G^2(QS) + G^2(MH) \end{aligned}$$

where $G^2(MH)$ is defined by the likelihood-ratio test of the difference between the S and QS models,

$$G^2(MH) \equiv G^2(S|QS) = G^2(S) - G^2(QS) . \quad (8.26) \quad \{\text{eq:mh}\}$$

The `gnm` package provides several model building convenience functions that facilitate fitting these and related models:

- `Diag(row, col, ...)` constructs a diagonals association factor for two (or more) factors with integer levels where the original factors are equal, and "." otherwise.
- `Symm(row, col, ...)` constructs an association factor giving equal levels to sets of symmetric cells. The QS model is specified using `Diag() + Symm()`.
- `Topo(row, col, ..., spec)` creates an association factor for two or more factors, as specified by an array of levels, which may be arbitrarily structured. Both `Diag()` and `Symm()` factors are special cases of `Topo()`.

The factor levels representing these association effects for a 4×4 table are shown below by their unique values in each array.

$$\text{Diag}_{4 \times 4} = \begin{bmatrix} 1 & . & . & . \\ . & 2 & . & . \\ . & . & 3 & . \\ . & . & . & 4 \end{bmatrix} \quad \text{Symm}_{4 \times 4} = \begin{bmatrix} 11 & 12 & 13 & 14 \\ 12 & 22 & 23 & 24 \\ 13 & 23 & 33 & 34 \\ 14 & 24 & 34 & 44 \end{bmatrix} \quad \text{Topo}_{4 \times 4} = \begin{bmatrix} 2 & 3 & 4 & 4 \\ 3 & 3 & 4 & 4 \\ 4 & 4 & 5 & 5 \\ 4 & 4 & 5 & 1 \end{bmatrix}$$

```
{ex:vision-glm}
```

EXAMPLE 8.7: Visual acuity

Example 4.13 presented the data on tests of visual acuity in the left and right eyes of a large sample of women working in the Royal Ordnance factories in World War II. A sieve diagram (Figure 4.8) showed that, as expected, most women had the same acuity in both eyes, but the off-diagonal cells had a pattern suggesting some form of symmetry.

The data set `VisualAcuity` contains data for both men and women in frequency form and for this example we subset this to include only the 4×4 table for women.

```
data("VisualAcuity", package="vcd")
women <- subset(VisualAcuity, gender=="female", select=-gender)
```

The four basic models of independence, quasi-independence, symmetry and quasi-symmetry for square tables are fit as shown below. We use `update()` to highlight the relations among these models in two pairs.

```
library(vcdExtra)
indep <- glm(Freq ~ right + left, data = women, family = poisson)
quasi <- update(indep, . ~ . + Diag(right, left))
```

```
symm <- glm(Freq ~ Symm(right, left), data = women, family = poisson)
qsymm <- update(symm, . ~ right + left + .)
```

The brief summary of goodness of fit of these models below shows that the QS model fits reasonably well, but none of the others do by likelihood-ratio tests or AIC or BIC.

```
models <- glmlist(indep, quasi, symm, qsymm)
vcdExtra::summarise(models)

## Model Summary:
##      LR Chisq Df Pr(>Chisq)  AIC  BIC
## indep    6672  9    0.0000 6654 6591
## quasi    199  5    0.0000  189  155
## symm     19  6    0.0038   7  -34
## qsymm     7  3    0.0638   1  -19
```

Beyond just saying that the QS model fits best, the reasons *why* it does can be seen in mosaic displays. Figure 8.11 compares the mosaics for the models of quasi-independence (accounting only for the diagonal cells) and quasi-symmetry (also accounting for symmetry). It can be seen in the left panel that the non-diagonal associations are largely symmetric, and also that when they differ, visual acuity in the two eyes are most likely to differ by only one eye grade.

```

labs <- c("High", "2", "3", "Low")
largs <- list(set_varnames = c(right="Right eye grade",
                                left="Left eye grade"),
              set_labels=list(right=labs, left=labs))
mosaic(quasi, ~right + left, residuals_type="rstandard",
       gp=shading_Friendly,
       labeling_args=largs,
       main="Quasi-Independence (women)")
mosaic(qsymm, ~right + left, residuals_type="rstandard",
       gp=shading_Friendly,
       labeling_args=largs,
       main="Quasi-Symmetry (women)")

```




Figure 8.11: Mosaic displays comparing the models of quasi-independence and quasi-symmetry for visual acuity in women. Fig: vision-mosaics

Finally, as usual, `anova()` can be used to carry out specific tests of nested models. For example, the test of marginal homogeneity Eqn. (8.26) compares models S and QS and shows here that the marginal probabilities for the left and right eyes differ.

```

anova(symm, qsymm, test="Chisq")

## Analysis of Deviance Table
##
## Model 1: Freq ~ Symm(right, left)
## Model 2: Freq ~ right + left + Symm(right, left)
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1         6      19.25
## 2         3       7.27  3      12    0.0075 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

△

{ex:hauser1}

EXAMPLE 8.8: Hauser's occupational mobility table

The data `Hauser79` in `vcdExtra`, from Hauser (1979), gives a 5×5 table in frequency form cross-classifying 19,912 individuals in the United States by father's occupation and son's first occupation. The occupational categories are represented by abbreviations, of Upper Non-Manual (UpNM), Lower Non-Manual (LoNM), Upper Manual (UpM), Lower Manual (LoM) and Farm. These data were also analysed by Powers and Xie (2008).

```
data("Hauser79", package="vcdExtra")
strutable(~Father+Son, data=Hauser79)
```

##	Son	UpNM	LoNM	UpM	LoM	Farm
## Father						
## UpNM		1414	521	302	643	40
## LoNM		724	524	254	703	48
## UpM		798	648	856	1676	108
## LoM		756	914	771	3325	237
## Farm		409	357	441	1611	1832

Before fitting any models, it is useful to calculate and plot the observed local log odds ratios, as we did in Example 8.5 to see the patterns in the data that need to be accounted for. These are calculated using `loddsratio()`.

```
hauser.tab <- xtabs(Freq ~ Father+Son, data=Hauser79)
(lor.hauser <- loddsratio(hauser.tab))
```

##	Son				
## Father	UpNM:LoNM	LoNM:UpM	UpM:LoM	LoM:Farm	
## UpNM:LoNM	0.6751	-0.1788	0.2623	0.09311	
## LoNM:UpM	0.1151	1.0025	-0.3461	-0.05788	
## UpM:LoM	0.3980	-0.4485	0.7896	0.10087	
## LoM:Farm	-0.3258	0.3815	-0.1660	2.76972	

This 4×4 table is graphed using `matplot()`, giving Figure 8.12.

```
matplot(as.matrix(lor.hauser), type='b', lwd=2,
        ylab='Local log odds ratio',
        xlab="Son's status comparisons",
        xaxt='n', cex.lab=1.2,
        xlim=c(1,4.5), ylim=c(-.5,3)
        )
abline(h=0, col='gray') # independence
abline(h=mean(lor.hauser$coefficients)) # mean
axis(side=1, at=1:4, labels=colnames(lor.hauser))
text(4, as.matrix(lor.hauser)[4,], rownames(lor.hauser),
     pos=4, col=1:4, xpd=TRUE, cex=1.2)
text(4, 3, "Father's status", cex=1.2)
```

Amongst the features here, you can see that there is a tendency for the odds ratio contrasting fathers in the non-manual categories (UpNM:LoNM) to decline with the adjacent comparisons of their sons' occupations. As well, the 2×2 table for fathers and sons in the LoM:Farm stands out as deserving some attention. These observed features will be smoothed by fitting models, as described below. For additional interpretation, you can always construct similar plots of the log odds ratios using the `fitted()` values from any of the models described below.

We begin by fitting the independence model and the quasi-independence model, where the diagonal parameters in the latter are specified as `Diag(Father, Son)`. As expected, given the




Figure 8.12: Plot of observed local log odds ratios in the Hauser79 data. The gray horizontal line at zero shows local independence; the black horizontal line shows the mean.

large frequencies in the diagonal cells, the quasi-independence model is a considerable improvement, but the fit is still very poor.

```
hauser.indep <- gnm(Freq ~ Father + Son, data=Hauser79, family=poisson)
hauser.quasi <- update(hauser.indep, ~ . + Diag(Father, Son))
vcdExtra::summarise(glmlist(hauser.indep, hauser.quasi))

## Model Summary:
##      LR Chisq Df Pr(>Chisq)  AIC  BIC
## hauser.indep    6170 16  0.00e+00 6138 6012
## hauser.quasi     683 11  2.01e-139 661  574
```

The pattern of associations can be seen in the mosaic displays for both models, shown in Figure 8.13.

```
mosaic(hauser.indep, ~Father+Son, main="Independence model",
       gp=shading_Friendly)
mosaic(hauser.quasi, ~Father+Son, main="Quasi-independence model",
       gp=shading_Friendly)
```

The mosaic for quasi-independence shows an approximately symmetric pattern of residuals, so we proceed to add `Symm(Father, Son)` to the model to specify quasi-symmetry.

```
hauser.qsymm <- update(hauser.indep,
                      ~ . + Diag(Father, Son) + Symm(Father, Son))
vcdExtra::summarise(hauser.qsymm)

## Model Summary:
##      LR Chisq Df Pr(>Chisq)  AIC  BIC
## hauser.qsymm    27.4  6  0.000119 15.4 -31.9
```

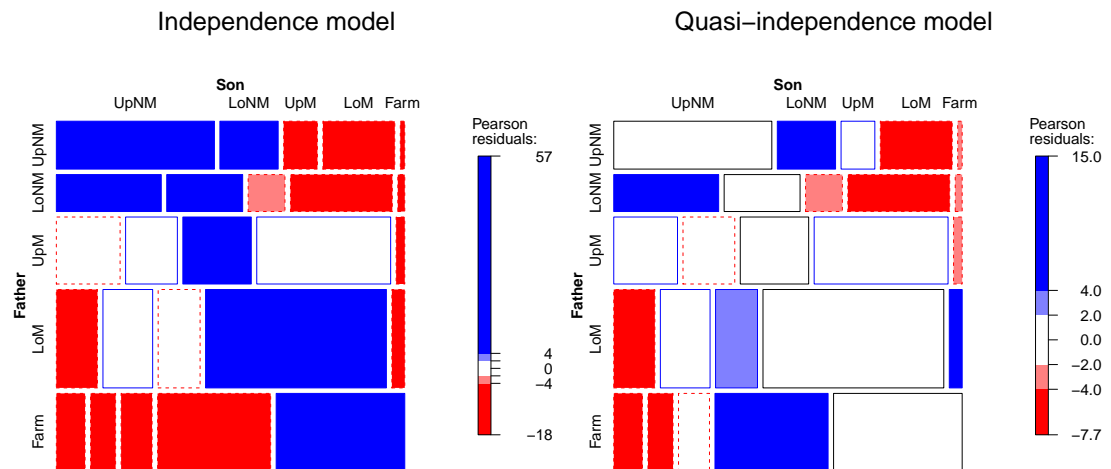


Figure 8.13: Mosaic displays for the Hauser79 data. Left: independence model; right: quasi-independence model. fig:hauser-mosaic1

This model represents a huge improvement in goodness of fit. With such a large sample size, it might be considered an acceptable fit. The remaining lack of fit is shown in the mosaic for this model, Figure 8.14.

```
mosaic(hauser.qsymm, ~Father+Son, main="Quasi-symmetry model",
       gp=shading_Friendly, residuals_type="rstandard")
```

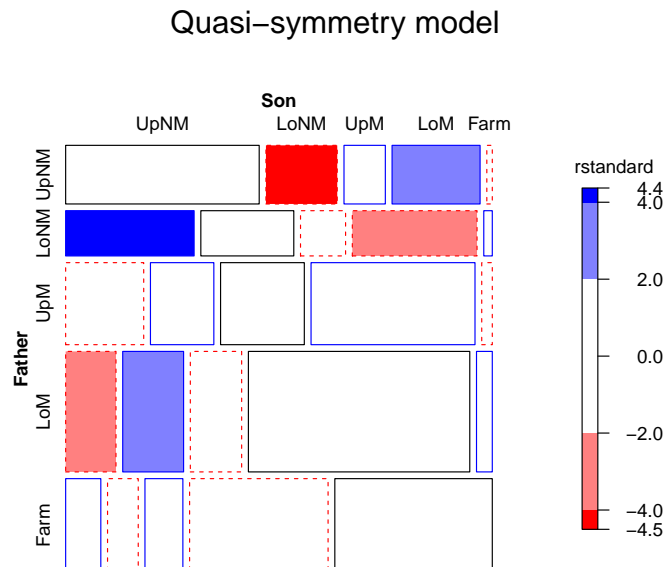


Figure 8.14: Mosaic display for the model of quasi-symmetry fit to the Hauser79 data. fig:hauser-mosaic2

The cells with the largest lack of symmetry (using standardized residuals) are those for the

upper and lower non-manual occupations, where the son of an upper manual worker is less likely to move to lower non-manual work than the reverse.

For cases like this involving structured associations in square tables, Hauser (1979) developed the more general idea of grouping the row and column categories into levels of an association factor based on similar values of residuals or local odds ratios observed from the independence model. Such models are called *topological models* or *levels models*, which are implemented in the `Topo()`.

To illustrate, Hauser suggested the following matrix of levels to account for the pattern of associations seen in Figure 8.13. The coding here takes the diagonal cell for the Farm category as the reference cell. Four other parameters are assigned by the numbers 2–5 to account for lack of independence.

```
levels <- matrix(c(
  2, 4, 5, 5, 5,
  3, 4, 5, 5, 5,
  5, 5, 5, 5, 5,
  5, 5, 5, 4, 4,
  5, 5, 5, 4, 1
), 5, 5, byrow=TRUE)
```

This model is fit using `Topo()` as shown below. It also provides a huge improvement over the independence model, with 4 additional parameters.

```
hauser.topo <- update(hauser.indep, ~ . + Topo(Father, Son, spec=levels))
vcdExtra::summarise(hauser.topo)

## Model Summary:
##           LR Chisq Df Pr(>Chisq)  AIC   BIC
## hauser.topo    66.6 12   1.4e-09 42.6 -52.2
```

As with other models fit using `gnm()`, you can extract the coefficients for particular terms using `pickCoef()`.

```
as.vector((coef(hauser.topo)[pickCoef(hauser.topo, "Topo")]))

## [1] -1.813 -2.497 -2.803 -3.403
```

The models fit in this example are summarized below. Note that AIC prefers the quasi-symmetry model, `hauser.quasi`, while, because of the large sample size, BIC prefers the topological model, `hauser.topo`.

```
vcdExtra::summarise(glmlist(hauser.indep, hauser.quasi, hauser.qsymm, hauser.topo))

## Model Summary:
##           LR Chisq Df Pr(>Chisq)  AIC   BIC
## hauser.indep    6170 16   0.000000 6138 6012
## hauser.quasi     683 11   0.000000 661 574
## hauser.qsymm      27 6   0.000119 15 -32
## hauser.topo      67 12   0.000000 43 -52
```

8.7.2 Ordinal square tables

{sec:sq-ord}

The theory presented in Section 8.7.1 treats the row and column variables as nominal. In many instances, such as Example 8.8, the variable categories are also ordered, yet these models do not exploit their ordinal nature. In such cases, the models such as uniform association ($L \times L$), row effects, RC and others discussed in Section 8.6 can be combined with terms for quasi-independence and symmetry of the remaining associations.

For example, the $L \times L$ model Eqn. (8.20) of uniform association applies directly to square tables, and, for square tables, can also be amended to include a diagonals term, `Diag()`, giving a model of *quasi-uniform association*. In this model, all adjacent 2×2 sub-tables not involving diagonal cells have a common local odds ratio.

A related model is the **crossings model** (Goodman, 1972). This hypothesizes that there are different difficulty parameters for crossing from one category to the next, and that the associations between categories decreases with their separation. In the crossings model for an $I \times I$ table, there are $I - 1$ crossings parameters, $\nu_1, \nu_2, \dots, \nu_{I-1}$. The association parameters, λ_{ij}^{AB} have the form of the product of the intervening ν parameters,

$$\lambda_{ij}^{AB} = \begin{cases} \prod_{k=j}^{i-1} \nu_k & : i > j \\ \prod_{k=i}^{j-1} \nu_k & : i < j \end{cases}$$

This model can also be cast in *quasi* form, by addition of a `Diag` term to fit the main diagonal cells. See Powers and Xie (2008, §4.4.7) for further details of this model. The `Crossings()` function in `vcdExtra` implements such crossings terms.

{ex:hauser2}

EXAMPLE 8.9: Hauser's occupational mobility table

Without much comment or detail, for reference we first fit some of the ordinal models to the `Hauser79` data: Uniform association ($L \times L$), row effects, and the RC(1) model.

```
Fscore <- as.numeric(Hauser79$Father) # numeric scores
Sscore <- as.numeric(Hauser79$Son)    # numeric scores

# uniform association
hauser.UA <- update(hauser.indep, ~ . + Fscore*Sscore)
# row effects model
hauser.roweff <- update(hauser.indep, ~ . + Father*Sscore)
# RC model
hauser.RC <- update(hauser.indep, ~ . + Mult(Father, Son), verbose=FALSE)
```

All of these fit very poorly, yet they are all substantial improvements over the independence model.

```
vcdExtra::summarise(glmlist(hauser.indep, hauser.UA,
                           hauser.roweff, hauser.RC))

## Model Summary:
##      LR Chisq Df Pr(>Chisq)  AIC  BIC
## hauser.indep    6170 16  0.0e+00 6138 6012
## hauser.UA       2281 15  0.0e+00 2251 2132
## hauser.roweff   2080 12  0.0e+00 2056 1961
## hauser.RC       685  9  9.3e-142  667  596
```

The $L \times L$ model, `hauser.UA` might be improved by ignoring the diagonals, and, indeed it is.

```
hauser.UAdiag <- update(hauser.UA, ~ . + Diag(Father, Son))
anova(hauser.UA, hauser.UAdiag, test="Chisq")

## Analysis of Deviance Table
##
## Model 1: Freq ~ Father + Son + Fscore + Sscore + Fscore:Sscore
## Model 2: Freq ~ Father + Son + Fscore + Sscore + Fscore:Sscore + Diag(Father,
##      Son)
##      Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1          15      2281
## 2          10         73 5      2208   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

In this model, the estimated common local log odds ratio—the coefficient for the linear-by-linear term `Fscore:Sscore` is

```
coef(hauser.UAdiag)[["Fscore:Sscore"]]

## [1] 0.1584
```

For comparisons not involving the diagonal cells, each step down the scale of occupational categories for the father multiplies the odds that the son will also be in one lower category by $\exp(0.158) = 1.172$, an increase of 17%.

The crossings model, with and without the diagonal cells can be fit as follows:

```
hauser.CR <- update(hauser.indep, ~ . + Crossings(Father, Son))
hauser.CRdiag <- update(hauser.CR, ~ . + Diag(Father, Son))
vcdExtra::summarise(glmlist(hauser.CR, hauser.CRdiag))

## Model Summary:
##      LR Chisq Df Pr(>Chisq)  AIC  BIC
## hauser.CR      89.9 12 1.00e-13 65.9 -28.9
## hauser.CRdiag   64.2 9 2.03e-10 46.2 -24.9
```

The quasi-crossings model `hauser.CRdiag` has a reasonable G^2 fit statistic, and its interpretation and lack of fit is worth exploring further. The crossings coefficients ν can be extracted as follows.

```
nu <- coef(hauser.CRdiag)[pickCoef(hauser.CRdiag, "Crossings")]
names(nu) <- gsub("Crossings(Father, Son)C", "nu", names(nu), fixed=TRUE)
nu

##      nu1      nu2      nu3      nu4
## -0.4227 -0.3877 -0.2750 -1.4024
```

They indicate the steps between adjacent categories in terms of the barriers for a son moving to a lower occupational category. The numerically largest gap separates the lower non-manual category from farming.

In contrast to the `UAdiag` model, the quasi-crossing model with diagonal terms implies that all 2×2 off-diagonal sub-tables are independent, i.e., the local odds ratios are all equal to 1.0.

The reasons for lack of fit of this model can be seen in the corresponding mosaic display, shown in Figure 8.15

```
mosaic(hauser.CRdiag, ~Father+Son,
       gp=shading_Friendly, residuals_type="rstandard",
       main="Crossings() + Diag()")
```

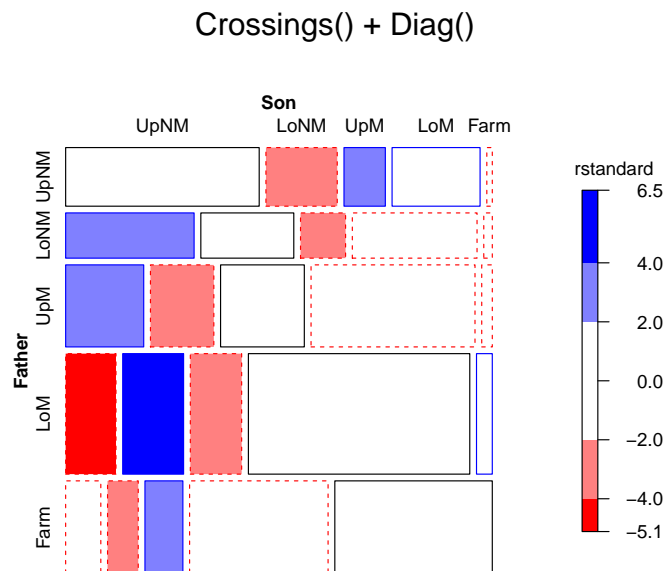


Figure 8.15: Mosaic display for the quasi-crossings model fit to the Hauser79 data. fig:hauser-mosaic3

It can be seen that lack of fit for this model is largely concentrated in the lower triangle, where the father's occupation is lower than that of his son.

In this example and the last, we have fit quite a few different models to the Hauser (1979) data. In presentations, articles and books it is common to summarize such a collection in a table, sorted by G^2 , degrees of freedom, AIC or BIC, to show their ordering along some metric. For instance, here we collect all the models fit in Example 8.8 and this example in a `glmlist()` and sort in decreasing order of AIC to show model fit by this measure.

```
modlist <- glmlist(hauser.indep, hauser.roweff, hauser.UA, hauser.UAdiag,
                  hauser.quasi, hauser.qsymm, hauser.topo,
                  hauser.RC, hauser.CR, hauser.CRdiag)
vcdExtra::summarise(modlist, sortby="AIC")

## Model Summary:
##      LR Chisq Df Pr(>Chisq)  AIC  BIC
## hauser.indep      6170 16  0.000000 6138 6012
## hauser.UA          2281 15  0.000000 2251 2132
## hauser.roweff      2080 12  0.000000 2056 1961
## hauser.RC           685  9  0.000000  667  596
## hauser.quasi        683 11  0.000000  661  574
## hauser.CR           90 12  0.000000   66  -29
## hauser.UAdiag       73 10  0.000000   53  -26
## hauser.CRdiag       64  9  0.000000   46  -25
## hauser.topo         67 12  0.000000   43  -52
## hauser.qsymm        27  6  0.000119   15  -32
```

When there are more than just a few models, a more useful display is a *model comparison plot* of measures like G^2/df , AIC or BIC against degrees of freedom. For example, Figure 8.16 plots BIC against Df in the result of `summarise()`. Because interest is focused on the smallest values of BIC and these values span a large range, BIC is shown on the log scale using `log="y"`; an offset is added to BIC to avoid negative values.

```
sumry <- vcdExtra::summarise(modlist)
mods <- substring(rownames(sumry), 8)
with(sumry, {
  plot(Df, BIC+55, cex=1.3, pch=19,
       xlab='Degrees of freedom', ylab='BIC (log scale)',
       log="y", cex.lab=1.2)
  pos <- ifelse(mods=="UAdiag", 1, 3)
  text(Df, BIC+55, mods, pos=pos, col='red', xpd=TRUE, , cex=1.2)
})
```




Figure 8.16: Model comparison plot for the models fit to the Hauser79 data fig:hauser-sumry-plot

Compared with the sorted tabular display shown above, such a plot sorts the models *both* by a measure of fit and by model complexity (degrees of freedom). Figure 8.16 shows that Hauser's topological model is best by BIC, but also shows that the next four best models by this measure are quite similar in terms of BIC. Similar plots for AIC and G^2/df show that the model of quasi-symmetry is favored by these measures.

△

8.8 Three-way and higher-dimensional tables

The models and methods for ordinal factors and square tables described in Section 8.6 and Section 8.7 extend readily to multidimensional tables with these properties for some of the factors.

In three-way tables, these models provide a more parsimonious account than the saturated model, $[ABC]$, and also allow simpler models than the general model of homogeneous association, $[AB][AC][BC]$ using scores for ordinal factors or terms for symmetry and diagonal factors in square layers.

For example, consider the case where all three factors are ordinal and the model of homogeneous association $[AB][AC][BC]$ fits poorly. In this case we can generalize the model of uniform association by assigning scores a , b and c and model the three-way association, λ_{ijk}^{ABC} as

$$\lambda_{ijk}^{ABC} = \gamma a_i b_j c_k$$

with only one more parameter. This gives the model of **uniform interaction** (or *homogeneous uniform association*)

$$\{eq:uni-inter\} \quad , \log m_{ijk} = \mu + \lambda_i^A + \lambda_j^B + \lambda_k^C + \lambda_{ij}^{AB} + \lambda_{ik}^{AC} + \lambda_{jk}^{BC} + \gamma a_i b_j c_k . \quad (8.27)$$

This model posits that (with equally spaced scores) all local odds ratios θ_{ijk} in adjacent rows, columns and layers are constant,

$$\log \theta_{ijk} = \gamma \quad \forall \quad i, j, k$$

The homogeneous association model is the special case of $\log \theta_{ijk} = \gamma = 0$.

A less restricted model of **heterogeneous uniform association** retains the linear-by-linear form of association for factors A and B , but allows the strength of this association to vary over layers, C , representing λ_{ijk}^{ABC} as

$$\lambda_{ijk}^{ABC} = (\gamma + \gamma_k) a_i b_j$$

with the constraint $\sum_k \gamma_k = 0$. This model is equivalent to fitting separate models of uniform association at each level k of factor C and gives estimates of the conditional local log odds ratios, $\log \theta_{ij(k)} = \gamma + \gamma_k$.

Following the development in Section 8.6 there is a large class of other models for ordinal factors (see Figure 8.6), where not all factors are assigned scores. For three-way tables, these can be represented in homogeneous form when the two-way association of A and B is the same for all levels of C , or in a heterogeneous form, when it varies over C .

Similarly, the models for square tables described in Section 8.7 extend to three-way tables with several layers (strata), allowing both homogeneous and heterogeneous terms for diagonals and symmetry describing the AB association over levels of C .

{ex:vision-glm2}

EXAMPLE 8.10: Visual acuity

We continue the analysis of the `VisualAcuity` data, but now consider the three-way, $4 \times 4 \times 2$ table comprising both men and women. The main questions here are whether the pattern of quasi-symmetry observed in the analysis for women also pertains to men and whether there is heterogeneity of the association between right, left acuity across gender.

A useful first step for n -dimensional tables is to consider the models composed of all 1-way, 2-way, \dots n -way terms as a quick overview. The function `Kway()` in `vcdExtra` package does this automatically, returning a "glmList" object containing the fitted models.⁹

⁹For completeness, this also fits the 0-way model, corresponding to $\log m_{ijk\dots} = \mu$, or the model formula `Freq ~ 1`.

```
vis.kway <- Kway(Freq ~ right + left + gender, data=VisualAcuity)
vcdExtra::summarise(vis.kway)
```

```
## Model Summary:
##      LR Chisq Df Pr(>Chisq)    AIC    BIC
## kway.0  13631 31  0.000000 13569 13344
## kway.1   9686 24  0.000000  9638  9463
## kway.2    28  9  0.000793   10   -55
## kway.3     0  0  0.000000    0     0
```

This shows that the model of homogeneous association `kway.2` ($[RL][RG][LG]$) does not fit well, but it doesn't account for diagonal agreement or symmetry to simplify the associations.

As a basis for comparison, we first fit the simple models of quasi-independence and quasi-symmetry that do not involve *gender*, asserting the same pattern of diagonal and off-diagonal cells for males and females.

```
vis.indep <- glm(Freq ~ right + left + gender, data = VisualAcuity,
                 family=poisson)
vis.quasi <- update(vis.indep, . ~ . + Diag(right, left))
vis.qsymm <- update(vis.indep, . ~ . + Diag(right, left) + Symm(right, left))

vis.models <- glmlist(vis.indep, vis.quasi, vis.qsymm)
vcdExtra::summarise(vis.models)
```

```
## Model Summary:
##      LR Chisq Df Pr(>Chisq)    AIC    BIC
## vis.indep   9686 24  0.00e+00 9638 9463
## vis.quasi    449 20  0.00e+00  409  263
## vis.qsymm    184 18  1.76e-29  148   17
```

The model of homogeneous quasi-symmetry fits quite badly, even worse than the all two-way association model. We can see why in the mosaic for this model, shown in Figure 8.17.

```
mosaic(vis.qsymm, ~ gender + right + left, condvars="gender",
        residuals_type="rstandard", gp=shading_Friendly,
        labeling_args=largs,
        main="Homogeneous quasi-symmetry")
```

It can be seen in Figure 8.17 that the pattern of residuals for men and women are nearly completely opposite in the upper and lower portions of the plot: men have positive residuals in the same right, left cells where women have negative residuals, and vice-versa. In particular, the diagonal cells of both tables have large absolute residuals, because the term `Diag(right, left)` fits a common set of diagonals for both men and women.

We can correct for this by allowing separate diagonal and symmetry terms, given as interactions of *gender* with `Diag()` and `Symm()`.

```
vis.hetdiag <- update(vis.indep, . ~ . + gender*Diag(right, left) +
                     Symm(right, left))
vis.hetqsymm <- update(vis.indep, . ~ . + gender*Diag(right, left) +
                     gender*Symm(right, left))
vis.hetmodels <- glmlist(vis.qsymm, vis.hetdiag, vis.hetqsymm)
vcdExtra::summarise(vis.hetmodels)

## Model Summary:
##      LR Chisq Df Pr(>Chisq)    AIC    BIC
```

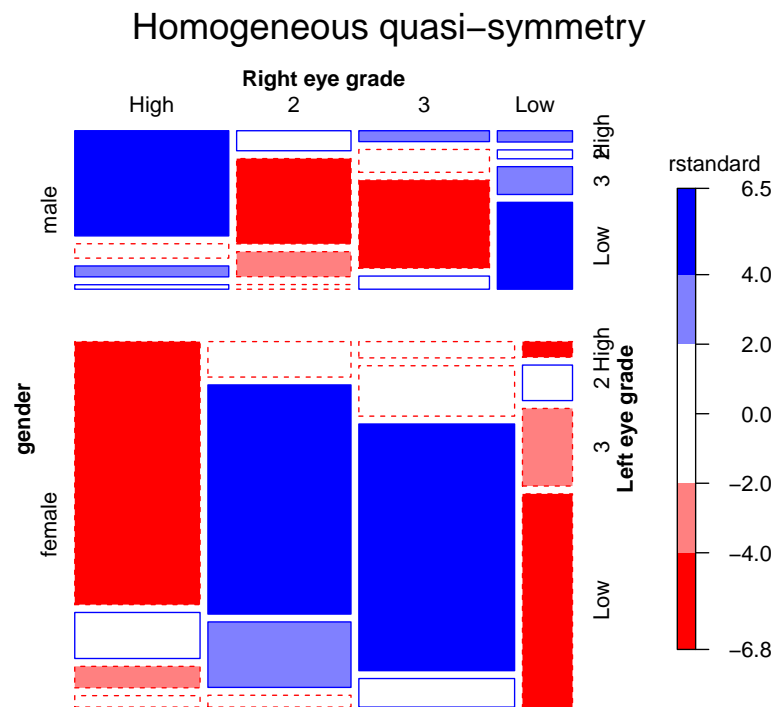


Figure 8.17: Mosaic display for the model of homogeneous quasi-symmetry fit to the VisualAcuity data.

```
## vis.qsymm      183.7 18      0.0000 147.7 16.7
## vis.hetdiag    52.3 14      0.0000 24.3 -77.6
## vis.hetqsymm   17.7 9      0.0384 -0.3 -65.8
```

Note that the model `vis.hetqsymm` fits better than the model `vis.hetdiag` in absolute terms and by AIC, but the latter, with fewer parameters, fits better by BIC. The mosaic for the model `vis.hetqsymm` is shown in Figure 8.18.

```
mosaic(vis.hetqsymm, ~ gender + right + left, condvars="gender",
        residuals_type="rstandard", gp=shading_Friendly,
        labeling_args=largs,
        main="Heterogeneous quasi-symmetry")
```

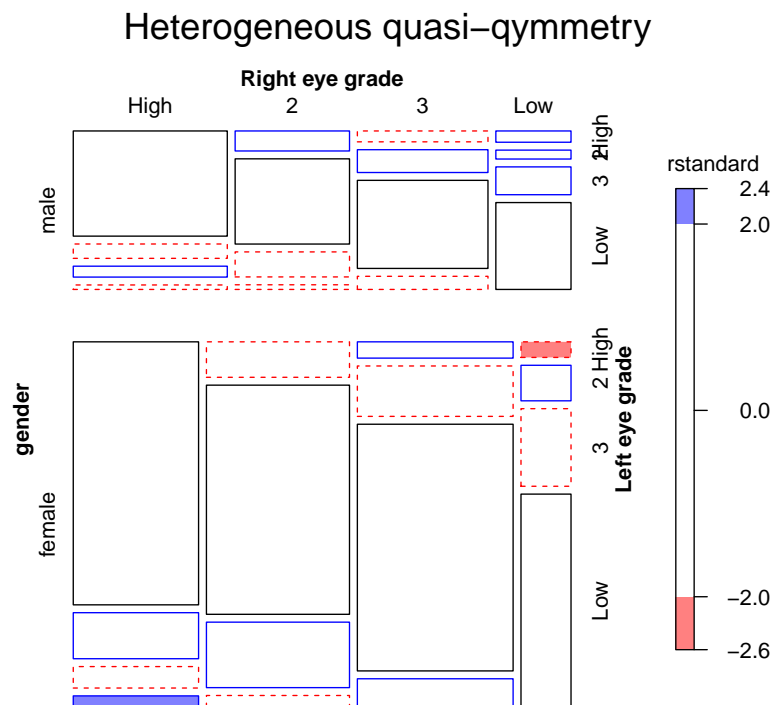


Figure 8.18: Mosaic display for the model of heterogeneous quasi-symmetry fit to the VisualAcuity data. `fig:vision2-hetqsymm`

As in the two-way case, this model now fits the diagonal cells in each table exactly, effectively ignoring this part of the association between right and left eye acuity. All remaining residuals are relatively small in magnitude, except for the two opposite off-diagonal cells (Low, High) and (High, Low) in the table for women.

The substantive interpretation of this example is that visual acuity is largely the same (diagonal cells) in the right and left eyes of both men and women. Ignoring the diagonal cells, when visual acuity differs, both men and women exhibit approximately symmetric associations. However, deviations from symmetry (Figure 8.17) are such that men are slightly more likely to have a lower grade in the right eye, while women are slightly more likely to have a higher grade in the right eye.

△

8.9 Influence and diagnostic plots for loglinear models

{sec:loglin

8.10 Multivariate responses

{sec:loglin

In many studies, there may be *several* categorical responses observed along with one or more explanatory variables. In a clinical trial, for example, the efficacy of a drug might be the primary response, but the occurrence of side-effects might give rise to additional response variables of substantive interest. Or, in a study of occupational health, the occurrence of two or more distinct symptoms might be treated as response variables.

If there are *no* explanatory variables, then the problem is simply to understand the joint distribution of the response categories, and the loglinear models and graphical displays described earlier are sufficient. Otherwise, in these cases we usually wish to understand how the various responses are affected by the explanatory variables. Moreover, it may also be important to understand how the association between the categorical responses depends on the explanatory variables. That is, we would like to study how *both* the marginal distributions of the responses, and their joint distribution depends on the predictors. In the occupational health example, the goal might be to understand both how the prevalence of several symptoms varies with one or more predictors, and how the association (loosely, “correlation”) among those symptoms varies with those predictors.

Although the general loglinear model is often used in these situations, there are special reparameterizations that may be used to separate the *marginal* dependence of each response on the explanatory variables from the relationship of the *association* among the responses on the explanatory variables.

Let us say that categorical responses, Y_1, Y_2, \dots have been observed, together with possible explanatory variables, X_1, X_2, \dots , and let $\pi_{ij\dots}$ be the joint probability of all the responses and explanatory variables; we also use \mathbf{x} to refer to the values of X_1, X_2, \dots .

Note that the minimal model of independence of all responses from each other and from the explanatory variables is the loglinear model $[Y_1][Y_2] \cdots [X_1 X_2 \cdots]$ (i.e., all associations among the X_i must be included). A no-effect model, in which the responses do not depend on the explanatory variables, but may be associated among themselves is $[Y_1 Y_2 \cdots][X_1 X_2 \cdots]$. However, these models do not separate the individual (marginal) effects of $X_1, X_2 \dots$ on each Y_i from their associative effects on the joint relationships among the Y_i .

There are three useful general approaches which *do* separate these effects:

1. Model the marginal dependence of each response, Y_i separately on X_1, X_2, \dots , and, in addition, model the interdependence among the responses, Y_1, Y_2, \dots .¹⁰
2. Model the joint dependence of all responses on X_1, X_2, \dots , but parameterized so that marginal and associative effects are delineated.
3. Construct simultaneous models, estimated together, for the marginal and joint dependence of the responses on the explanatory variables.

The first approach is the simplest, an informative starting place, and is satisfactory in the (often unlikely) case that the responses are not associated, or if the associations among responses

¹⁰For quantitative responses, this is roughly analogous to fitting univariate response models for each Y_i , followed by something like a principal component analysis of the relationships among the Y_i . But in this case, the multivariate linear model, $\mathbf{Y} = \mathbf{XB} + \mathbf{E}$ provides a general solution.

do not vary much over the explanatory variables (i.e., no terms like $[Y_1 Y_2 X_j]$ are required). In the clinical trial example, we would construct separate loglinear or logit models for efficacy of the drug, and for occurrence of side-effects, and supplement these analyses with mosaic or other displays showing the relations between efficacy and side-effects and a model for their joint association. If those who improve with the drug also show more serious side effects, the worth of the treatment would be questioned. A limitation of this method is that it does not provide an overall model comprising these effects.

In the second approach, the joint probabilities, $\pi_{ij\dots}$, are recast to give separate information regarding the dependence of the univariate marginal probabilities $\pi_{i\bullet}, \pi_{\bullet j}, \dots$, on the explanatory variables and the dependence of the intra-response associations on the explanatory variables. The VGAM package provides several versions of this approach with the function `vglm()` (for *vector generalized linear model*).

The third approach, developed, for example, by Lang and Agresti (1994), is the most general, and provides a scheme to represent a model $\mathcal{J}(\bullet)$ for the joint distributions of the X, Y variables together with a model $\mathcal{M}(\bullet)$ for their first-order marginal distributions. The joint models are typically loglinear models, ranging from the mutual independence model, $\mathcal{J}(I) = [Y_1][Y_2][\dots][X_1][X_2][\dots]$ to the saturated model, $\mathcal{J}(S) = [Y_1 Y_2 \dots X_1 X_2 \dots]$, while the marginal models are logit models for the response variables. The combined model, denoted $\mathcal{J}(\bullet) \cap \mathcal{M}(\bullet)$, is estimated simultaneously by maximum likelihood. This approach is implemented in R in the hmmm package (hierarchical multinomial marginal models). However, model specification in this implementation is complicated, and it will not be considered further here.

8.10.1 Bivariate, binary response models

We focus here on two related models reflecting the second approach, as discussed by McCullagh and Nelder (1989, Section 6.5). We consider here only the case of two binary responses, though the general approach can be applied to $R > 2$ responses Y_1, Y_2, \dots, Y_R , and these may be polytomous or ordinal.

Let \mathbf{x} refer to the values of all the explanatory variables and let $\pi_{ij}(\mathbf{x})$ be the joint probabilities in cell $Y_1 = i, Y_2 = j$. The essential idea of the *bivariate logistic model* arises from a linear transformation of the cell probabilities $\boldsymbol{\pi}$ to interpretable functions of the marginal probabilities (logits) and their association (odds ratio), a mapping of $\boldsymbol{\pi} \rightarrow \boldsymbol{\eta}$,

$$\begin{aligned}\eta_1 &= \text{logit}(\pi_{1\bullet}) \\ \eta_2 &= \text{logit}(\pi_{\bullet 1}) \\ \eta_{12} &= \frac{\pi_{11} \pi_{22}}{\pi_{12} \pi_{21}}\end{aligned}\tag{8.28} \quad \{\text{eq:blogits}\}$$

The predictors in \mathbf{x} are then taken into account by considering models that relate $\boldsymbol{\pi}$ to \mathbf{x} through $\boldsymbol{\eta}$,

$$\begin{aligned}\eta_1 &= \mathbf{x}_1^\top \boldsymbol{\beta}_1 \\ \eta_2 &= \mathbf{x}_2^\top \boldsymbol{\beta}_2 \\ \eta_{12} &= \mathbf{x}_{12}^\top \boldsymbol{\beta}_{12}\end{aligned}\tag{8.29} \quad \{\text{eq:blogits2}\}$$

where $\mathbf{x}_1, \mathbf{x}_2$ and \mathbf{x}_{12} are subsets of the predictors in \mathbf{x} for each sub-model, and $\boldsymbol{\beta}_1, \boldsymbol{\beta}_2$ and $\boldsymbol{\beta}_{12}$ are the corresponding parameters to be estimated.

McCullagh and Nelder (1989) arrive at this joint bivariate model in two steps. First, transform the cell probabilities π to a vector of probabilities γ which also includes the univariate margins, given by

$$\gamma = L\pi \quad (8.30) \quad \{\text{eq:gamma1}\}$$

where L is a matrix of 0s and 1s of the form of a factorial design matrix. In the 2×2 case,

$$\{\text{eq:gamma2}\} \quad \gamma = \begin{pmatrix} \pi_{1\bullet} \\ \pi_{2\bullet} \\ \pi_{\bullet 1} \\ \pi_{\bullet 2} \\ \pi_{11} \\ \pi_{12} \\ \pi_{21} \\ \pi_{22} \end{pmatrix} = \begin{bmatrix} 1 & 1 & 0 & 0 \\ 0 & 0 & 1 & 1 \\ 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{pmatrix} \pi_{11} \\ \pi_{12} \\ \pi_{21} \\ \pi_{22} \end{pmatrix}. \quad (8.31)$$

There are of course only three linearly independent probabilities, because $\sum \sum \pi_{ij} = 1$. In the second step, the bivariate logistic model is formulated in terms of factorial contrasts on the elements of γ which express separate models for the two logits and the log odds. The model is expressed as

$$\{\text{eq:eta1}\} \quad \eta = C \log \gamma = C \log L\pi, \quad (8.32)$$

where C is a matrix of contrasts. In the 2×2 case, the usual contrasts may be defined by

$$\{\text{eq:eta2}\} \quad \eta = \begin{pmatrix} \eta_1 \\ \eta_2 \\ \eta_{12} \end{pmatrix} = \begin{pmatrix} \text{logit } \pi_{1\bullet} \\ \text{logit } \pi_{\bullet 1} \\ \theta_{12} \end{pmatrix} = \begin{bmatrix} 1 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & -1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & -1 & -1 & 1 \end{bmatrix} \begin{pmatrix} \pi_{1\bullet} \\ \pi_{2\bullet} \\ \pi_{\bullet 1} \\ \pi_{\bullet 2} \\ \pi_{11} \\ \pi_{12} \\ \pi_{21} \\ \pi_{22} \end{pmatrix} \quad (8.33)$$

Thus, we are modeling the marginal odds of each response, together with the log odds ratio θ_{12} simultaneously.

Specific models are then formulated for the dependence of $\eta_1(x)$, $\eta_2(x)$ and $\eta_{12}(x)$ on some or all of the the explanatory variables. For example, with one quantitative explanatory variable, x , the model

$$\{\text{eq:bi-logit1}\} \quad \begin{pmatrix} \eta_1 \\ \eta_2 \\ \eta_{12} \end{pmatrix} = \begin{pmatrix} \alpha_1 + \beta_1 x \\ \alpha_2 + \beta_2 x \\ \theta \end{pmatrix} \quad (8.34)$$

asserts that the log odds of each response changes linearly with x , while the odds ratio between the responses remains constant. In the general form given by McCullagh and Nelder (1989) the submodels in Eqn. (8.34) may each depend on the explanatory variables in different ways. For example, the logits could both depend quadratically on x , while an intercept-only model could be posited for the log odds ratio.

The second model is the **bivariate loglinear model**, the special case obtained by taking $L = I$ in Eqn. (8.30) and Eqn. (8.32) so that $\gamma = \pi$. Then a loglinear model of the form

$$\eta(x) = C \log \pi$$

expresses contrasts among log probabilities as linear functions of the explanatory variables. For the 2×2 case, we take the contrasts C as shown below

{eq:eta3}

$$\eta = \begin{pmatrix} l_1 \\ l_2 \\ \eta_{12} \end{pmatrix} = \begin{bmatrix} 1 & 1 & -1 & -1 \\ 1 & -1 & 1 & -1 \\ 1 & -1 & 1 & -1 \end{bmatrix} \begin{pmatrix} \log \pi_{11} \\ \log \pi_{12} \\ \log \pi_{21} \\ \log \pi_{22} \end{pmatrix} \quad (8.35)$$

and models for the dependence of $l_1(x)$, $l_2(x)$ and $\eta_{12}(x)$ are expressed in the same way as in Eqn. (8.34). The estimates of the odds ratio, η_{12} are the same under both models. The marginal functions are parameterized differently, however, but lead to similar predicted probabilities.

In R, bivariate logistic models of the form Eqn. (8.28) and Eqn. (8.29) can be fit using `vglm()` with the `binom2.or()` family in the **VGAM** package.¹¹ The fitting and graphing of these models is illustrated in the next example.

{ex:coalminers}

EXAMPLE 8.11: Breathlessness and wheeze in coal miners

In Example 4.11 we examined the association between the occurrence of two pulmonary conditions, breathlessness and wheeze, among coal miners classified by age (Ashford and Sowden, 1970). Figure 4.5 showed fourfold displays focused on the odds ratio for the co-occurrence of these symptoms, and Figure 4.6 plotted these odds ratios against age directly. Here, we consider models which examine the changes in prevalence of the two symptoms over age, together with the changes in their association.

Plotting bivariate response data

As a starting point and overview of what is necessary for bivariate response models, we calculate the empirical log odds for breathlessness and for wheeze, and the log odds ratio for their association in each 2×2 table. The log odds ratios are the same values plotted in Figure 4.6 (but the youngest age group was not included in the earlier analysis).

The `CoalMiners` data is $2 \times 2 \times 9$ table. For convenience in this analysis (and for use with **VGAM**) we convert it to a 4×9 data frame, and relabel the columns to use the combinations of ("B", "b") and ("W", "w") to represent the conditions of breathlessness and wheeze, where the upper case letter indicates presence of the condition. A variable `age` is also created, using the midpoints of the age categories.

```
data("CoalMiners", package="vcd")
coalminers <- data.frame(t(matrix(aperm(CoalMiners, c(2,1,3)),
                                   4, 9)))
colnames(coalminers) <- c("BW", "Bw", "bW", "bw")
coalminers$age <- c(22, 27, 32, 37, 42, 47, 52, 57, 62)
coalminers
```

##	BW	Bw	bW	bw	age
## 1	9	7	95	1841	22
## 2	23	9	105	1654	27
## 3	54	19	177	1863	32
## 4	121	48	257	2357	37
## 5	169	54	273	1778	42
## 6	269	88	324	1712	47
## 7	404	117	245	1324	52
## 8	406	152	225	967	57
## 9	372	106	132	526	62

¹¹This package also provides for bivariate and trivariate loglinear models with `loglinb2()` and `loglinb2`.

With the data in this form, the simple function `blogits()` calculates the logits and log odds ratios corresponding to Eqn. (8.28). The `add` argument accommodates cases where there are very small, or 0 frequencies in some cells, and it is common to add a small constant, such as 0.5 to each cell in calculating *empirical logits*.

```
# calculate bivariate logits and OR
blogits <- function(Y, add=0) {
  Y <- Y + add
  L <- matrix(0, nrow(Y), 3)
  L[,1] <- log( (Y[,1] + Y[,2]) / (Y[,3] + Y[,4]) )
  L[,2] <- log( (Y[,1] + Y[,3]) / (Y[,2] + Y[,4]) )
  L[,3] <- log( (Y[,1] * Y[,4]) / ((Y[,2] * Y[,3])) )
  colnames(L) <- c("logit1", "logit2", "logOR")
  L
}
```

This function is then used to calculate the empirical logits and log odds as follows:

```
logitsCM <- blogits(coalminers[,1:4], add=0.5)
colnames(logitsCM)[1:2] <- c("logitB", "logitW")
logitsCM

##          logitB  logitW logOR
## [1,] -4.7357 -2.8684 3.196
## [2,] -3.9766 -2.5572 3.658
## [3,] -3.3171 -2.0939 3.379
## [4,] -2.7332 -1.8482 3.133
## [5,] -2.2149 -1.4201 3.007
## [6,] -1.7387 -1.1092 2.777
## [7,] -1.1012 -0.7968 2.922
## [8,] -0.7581 -0.5722 2.437
## [9,] -0.3190 -0.2259 2.632
```

We plot these as shown below, using `matplot()`, which is convenient for plotting multiple columns against a given horizontal variable, age here.¹² For ease of interpretation of the log odds, we also use right vertical axis showing the equivalent probabilities for breathlessness and wheeze.

```
col <- c("blue", "red", "black")
pch <- c(15, 17, 16)
age <- coalminers$age

op <- par(mar=c(4, 4, 1, 4)+.2)
matplot(age, logitsCM, type="p",
  col=col, pch=pch, cex=1.2, cex.lab=1.25,
  xlab="Age", ylab="Log Odds or Odds Ratio")
abline(lm(logitsCM[,1] ~ age), col=col[1], lwd=2)
abline(lm(logitsCM[,2] ~ age), col=col[2], lwd=2)
abline(lm(logitsCM[,3] ~ age), col=col[3], lwd=2)

# right probability axis
probs <- c(.01, .05, .10, .25, .5)
axis(4, at=qlogis(probs), labels=probs)
mtext("Probability", side=4, cex=1.2, at=-2, line=2.5)
# curve labels
```

¹²It is actually a small graphical misdemeanor to plot logits and odds ratios on the same vertical axis because they are not strictly commensurable. We plead guilty with the explanation that this graph shows what we want to see here and does not distort the data.

```
text(age[2], logitsCM[2,1]+.5, "Breathlessness", col=col[1], pos=NULL, cex=1.2)
text(age[2], logitsCM[2,2]+.5, "Wheeze", col=col[2], pos=NULL, cex=1.2)
text(age[2], logitsCM[2,3]-.5, "log OR\n(B|W)/(B|w)", col=col[3], pos=1, cex=1.2)
par(op)
```




Figure 8.19: Empirical logits and log odds ratio for breathlessness and wheeze in the CoalMiners data. The lines show separate linear regressions for each function. The right vertical axis shows equivalent probabilities for the logits.

In Figure 8.19 we see that both symptoms, while quite rare among young miners, increase steadily with age (or years working in the mine). By age 60, the probability is nearly 0.5 of having either condition. There is a hint of curvilinearity, particularly in the logit for breathlessness. The decline in the odds ratio with age may reflect selection, as miners who had retired for health or other reasons were excluded from the study.

Fitting `glm` models

Next, we illustrate what can easily be achieved using the standard `glm()` approach for loglinear models and why the bivariate models we described are more useful in this situation. `glm()` requires a data frame as input, so first reshape `CoalMiners` to a frequency data frame. For convenience, we simplify the variable names to `B` and `W`.

```
CM <- as.data.frame(CoalMiners)
colnames(CM)[1:2] <- c("B", "W")
str(CM)

## 'data.frame': 36 obs. of 4 variables:
```

```
## $ B      : Factor w/ 2 levels "B","NoB": 1 2 1 2 1 2 1 2 1 2 ...
## $ W      : Factor w/ 2 levels "W","NoW": 1 1 2 2 1 1 2 2 1 1 ...
## $ Age     : Factor w/ 9 levels "20-24","25-29",...: 1 1 1 1 2 2 2 2 3 3 ...
## $ Freq: num  9 95 7 1841 23 ...
```

As a point of comparison, we fit the mutual independence model, $[B][W][Age]$ and the baseline model for associated responses, $[BW][Age]$ which asserts that the association between B and W is independent of Age.

```
cm.glm0 <- glm(Freq ~ B + W + Age, data=CM, family=poisson)
cm.glm1 <- glm(Freq ~ B * W + Age, data=CM, family=poisson)
vcdExtra::summarise(glm1list(cm.glm0, cm.glm1))
```

```
## Model Summary:
##          LR Chisq Df Pr(>Chisq)   AIC   BIC
## cm.glm0      6939 25          0 6889 6694
## cm.glm1      2702 24          0 2654 2466
```

The baseline model `cm.glm1` fits very badly. We can see the pattern of the residual association in a mosaic display for this model shown in Figure 8.20. The formula argument here specifies the order of the variables in the mosaic.

```
vnames <- list(set_varnames = c(B="Breathlessness", W="Wheeze"))
lnames <- list(B=c("B", "b"), W = c("W", "w"))
mosaic(cm.glm1, ~ Age + B + W,
        labeling_args=vnames, set_labels=lnames)
```




Figure 8.20: Mosaic display for the baseline model, $[BW][Age]$, fit to the CoalMiners data fig:cm-mosaic1

As structured here, it is easy to see the increase in the prevalence of breathlessness and wheeze with age and the changing pattern of their association with age.

From Figure 8.19 and Figure 8.20, it is apparent that both breathlessness and wheeze increase with age, so we can model this by adding terms $[B \text{ Age}][W \text{ Age}]$ to the baseline model. This is the no-three-way interaction model, which could also be specified as $\text{Freq} \sim (B + W + \text{Age})^2$.

```
cm.glm2 <- glm(Freq ~ B * W + (B + W) * Age, data=CM, family=poisson)
vcdExtra::summarise(glmlist(cm.glm1, cm.glm2))

## Model Summary:
##      LR Chisq Df Pr(>Chisq)  AIC  BIC
## cm.glm1    2702 24  0.000000 2654 2466
## cm.glm2      27  8  0.000799   11  -52
```

The improvement in fit is substantial, and all terms are highly significant, yet, the residual $G^2(8)$ indicates there is still lack of fit.

```
library(car)
Anova(cm.glm2)

## Analysis of Deviance Table (Type II tests)
##
## Response: Freq
##      LR Chisq Df Pr(>Chisq)
## B      11026  1  <2e-16 ***
## W       7038  1  <2e-16 ***
## Age      887  8  <2e-16 ***
## B:W      3025  1  <2e-16 ***
## B:Age    1130  8  <2e-16 ***
## W:Age     333  8  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

One way to improve the model using the `glm()` framework is to make use of *Age* as a quantitative variable and add a term to allow the odds ratio for the $[BW]$ association to vary linearly with age. Here, we construct the variable *ageOR* using the midpoints of the *Age* intervals.

```
CM$age <- rep(seq(22, 62, 5), each=4)
```

In the `glm()` approach, the odds ratio cannot be modeled directly, but we can use the following trick: For each 2×2 subtable, the odds ratio can be parameterized in terms of the frequency in any one cell, say, n_{11k} , given that the marginal total n_{++k} is included in the model. We do this by adding a new interaction variable, *ageOR* having the value of *age* for the $(1, 1, k)$ cells and 0 otherwise.

```
CM$ageOR <- (CM$B=="B") * (CM$W=="W") * CM$age
cm.glm3 <- update(cm.glm2, . ~ . + ageOR)
vcdExtra::summarise(glmlist(cm.glm0, cm.glm1, cm.glm2, cm.glm3))

## Model Summary:
##      LR Chisq Df Pr(>Chisq)  AIC  BIC
## cm.glm0    6939 25  0.000 6889 6694
## cm.glm1    2702 24  0.000 2654 2466
## cm.glm2      27  8  0.001   11  -52
## cm.glm3       7  7  0.450   -7  -62
```

The model `cm.glm3`, with one more parameter, now fits reasonably well, having residual $G^2(7) = 6.80$. The likelihood ratio test of model `cm.glm3` against `cm.glm2`, which assumes equal odds

ratios over age, can be regarded as a test of the hypothesis of homogeneity of odds ratios, against the alternative that the [BW] association changes linearly with age. The `glm()` models fit in this example are summarized in Table 8.3.

{tab:ashmod}

Table 8.3: Loglinear models `glm()` fit to Ashford & Sowden data

Model	Terms	LR Chisq	G^2	Pr(>Chisq)	AIC	BIC
cm.glm0	[B][W][A]	6939.07	25	0.0000	6889.07	6693.73
cm.glm1	[BW][A]	2701.94	24	0.0000	2653.94	2466.42
cm.glm2	[BW][BA][WA]	26.69	8	0.0008	10.69	-51.82
cm.glm3	[BWx][BA][WA]	6.80	7	0.4498	-7.20	-61.89

```
anova(cm.glm2, cm.glm3, test="Chisq")

## Analysis of Deviance Table
##
## Model 1: Freq ~ B * W + (B + W) * Age
## Model 2: Freq ~ B + W + Age + ageOR + B:W + B:Age + W:Age
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1         8      26.7
## 2         7       6.8  1      19.9  8.2e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

This analysis, while useful, also shows the limitations of the `glm()` approach: (a) It doesn't easily allow us to represent and test the substantively interesting hypotheses regarding *how* the prevalence of the binary responses, B and W vary with Age, such as seen in Figure 8.19. (b) It doesn't represent the odds ratio for the [BW] association directly, but only through the coding trick we used here. Thus, it is difficult to interpret the coefficient for `ageOR` = -0.02613 in a substantively meaningful way, except that it shows that the odds ratio is decreasing.¹³

Fitting `vglm` models

The `vglm()` function in the **VGAM** package provides a very general implementation of these and other models for discrete multivariate responses. The family function, `binom2.or()` for binary logistic models allows some or all of the logits or odds ratio submodels to be constrained to be intercept-only (e.g., as in Eqn. (8.34)) and the two marginal distributions can be constrained to be equal.

Quantitative predictors (such as age, here), can be modeled linearly or nonlinearly, using `poly()` for a parametric fit, or smooth regression splines, as provided by the functions `ns()`, `bs()` and others in model formulas. In this illustration, we fit bivariate linear and quadratic models in age.

`vglm()` takes its input data in the wide form we called `coalminers` at the beginning of this example. We could use the 9-level factor, `Age` as we did with `glm()`, but we plan to use `age` as a numeric variable in all three submodels. The coefficients in these models will be more easily interpreted if we center age and express it as `agec` in units of five years, as shown below.

¹³Actually, the interpretability of the coefficient for the log odds ratio can be enhanced here by centering age, and representing its units in steps of 5 years, as we do below.

```
coalminers <- transform(coalminers, agec=(age-42)/5)
coalminers$Age <- dimnames(CoalMiners)[[3]]
coalminers
```

```
##      BW  Bw  bW   bw age agec   Age
## 1    9   7  95 1841  22   -4 20-24
## 2   23   9 105 1654  27   -3 25-29
## 3   54  19 177 1863  32   -2 30-34
## 4  121  48 257 2357  37   -1 35-39
## 5  169  54 273 1778  42    0 40-44
## 6  269  88 324 1712  47    1 45-49
## 7  404 117 245 1324  52    2 50-54
## 8  406 152 225  967  57    3 55-59
## 9  372 106 132  526  62    4 60-64
```

`vglm()` takes the 2×2 response frequencies as a 4-column matrix on the right hand side of the model formula. However, denoting the responses of failure and success by 0 and 1 respectively, it takes these in the order $y_{00}, y_{01}, y_{10}, y_{11}$. We specify the order below so that the logits are calculated for the occurrence of breathlessness or wheeze, rather than their absence.

```
library(VGAM)
#               00  01  10  11
cm.vglm1 <- vglm(cbind(bw, bW, Bw, BW) ~ agec,
                 binom2.or(zero=NULL), data=coalminers)
cm.vglm1

## Call:
## vglm(formula = cbind(bw, bW, Bw, BW) ~ agec, family = binom2.or(zero = NULL),
##      data = coalminers)
##
## Coefficients:
## (Intercept):1 (Intercept):2 (Intercept):3      agec:1
##      -2.2625      -1.4878          3.0219      0.5145
##      agec:2      agec:3
##      0.3254      -0.1314
##
## Degrees of Freedom: 27 Total; 21 Residual
## Residual deviance: 30.39
## Log-likelihood: -100.5
```

In this call, the argument `zero=NULL` indicates that none of the linear predictors, $\eta_1, \eta_2, \eta_{12}$ are modeled as constants.¹⁴

At this writing, there is no `anova()` method for the "vgam" objects produced by `vglm()`, but we can test the residual deviance of the model (against the saturated model) as follows, showing that this model has an acceptable fit.

```
(G2 <- deviance(cm.vglm1))

## [1] 30.39

# test residual deviance
1-pchisq(deviance(cm.vglm1), cm.vglm1@df.residual)

## [1] 0.08435
```

¹⁴The default, `zero=3` gives the model shown in Eqn. (8.34), with the odds ratio constant.

The estimated coefficients in this model are usefully shown as below, using the argument `matrix=TRUE` in `coef()`. Using `exp()` on the result gives values of odds that can be easily interpreted:

```
coef(cm.vglm1, matrix=TRUE)

##           logit(mu1) logit(mu2) log(oratio)
## (Intercept)   -2.2625   -1.4878     3.0219
## agec           0.5145     0.3254    -0.1314

exp(coef(cm.vglm1, matrix=TRUE))

##           logit(mu1) logit(mu2) log(oratio)
## (Intercept)     0.1041     0.2259    20.5304
## agec           1.6728     1.3846     0.8769
```

Thus, the odds of a miner showing breathlessness are multiplied by 1.67, a 67% increase, for each 5 years increase in age; similarly, the odds of wheeze are multiplied by 1.38, a 38% increase. The odds ratio for the association between the two symptoms are multiplied by 0.88, a 12% decrease over each 5 year interval.

The VGAM package has no special plot methods for "vglm" objects, but it is not hard to construct these using the methods we showed earlier in this example. First, we can obtain the fitted probabilities for the 4 response combinations using `fitted()` and the corresponding observed probabilities using `devar()`.

```
age <- coalminers$age
P <- fitted(cm.vglm1)
colnames(P) <- c("bw", "bW", "Bw", "BW")
head(P)

##           bw          bW          Bw          BW
## 1 0.9375 0.04941 0.004636 0.008483
## 2 0.9146 0.06364 0.006976 0.014778
## 3 0.8841 0.08003 0.010496 0.025368
## 4 0.8439 0.09748 0.015814 0.042767
## 5 0.7919 0.11384 0.023860 0.070420
## 6 0.7258 0.12591 0.035968 0.112337

Y <- devar(cm.vglm1)
```

In the left panel of Figure 8.21, we plot the fitted probabilities in the matrix `P` using `matplot()` and the observed probabilities in `Y` using `matpoints()`.

```
col <- c("red", "blue", "red", "blue")
pch <- c(1, 2, 16, 17)

op <- par(mar=c(5, 4, 1, 1)+.1)
matplot(age, P, type="l",
        col=col,
        lwd=2, cex=1.2, cex.lab=1.2,
        xlab="Age", ylab="Probability",
        xlim=c(20, 65))
matpoints(age, Y,
         pch=pch, cex=1.2, col=col)
# legend
text(64, P[9,]+c(0,.01,-.01,0), labels=colnames(P), col=col, cex=1.2)
```




Figure 8.21: Observed and fitted values for the combinations of breathlessness and wheeze in the binary logistic regression model `cm.vglm1`. Left: probabilities; right: on the log odds scale.

```
text(20, P[1,]+ c(0,.01, -.01, .01), labels=colnames(P), col=col, cex=1.2)
par(op)
```

The right panel of Figure 8.21 shows these on the log odds scale, produced using the same code as above, applied to the probabilities transformed using `qlogis()`, the quantile function for the logistic distribution.

```
lP <- qlogis(P)
lY <- qlogis(Y)
```

In Figure 8.19 we plotted the empirical logits and log odds using the function `blogits()` to transform frequencies to these values. An essentially identical plot can be produced by transforming the fitted and observed probabilities, as calculated below.

```
# blogits, but for B and W
logitsP <- blogits(P[,4:1])
logitsY <- blogits(Y[,4:1])
```

To test for nonlinearity in the prevalence of the symptoms or their odds ratio with age, we can fit a similar model using `poly()` or a smoothing spline, such as `ns()`. We illustrate this here using a bivariate model allowing quadratic effects of age on all three components.

```
cm.vglm2 <- vglm(cbind(bw, bW, Bw, BW) ~ poly(agec,2),
  binom2.or(zero=NULL), data=coalminers)
```

This model has a residual $G^2 = 16.963$ with 18 df. Compared to the linear model `cm.vglm1`, this represents a significant improvement in goodness of fit.

```
(LR <- deviance(cm.vglm1) - deviance(cm.vglm2))
```

```
## [1] 13.43
```



```
1 - pchisq(LR, cm.vglm1@df.residual - cm.vglm2@df.residual)

## [1] 0.003792
```

A plot of the fitted logits and log odds ratios under this model is shown in Figure 8.22. You can interpret this plot as showing that the statistical evidence for the quadratic model indicates some slight tendency for the prevalence of breathlessness and wheeze levels off slightly with age, particularly the former.




Figure 8.22: Observed (points) and fitted (lines) logits and log odds ratios for the quadratic binary logistic regression model `cm.vglm2`.

`fig:cm-vglm2-blogit}`

△

8.10.2 More complex models

When there is more than one explanatory variable and several responses, the methods described above using `glm()` and `vglm()` still apply. However, it is useful to begin with a more thorough visual examination of the relations within and between these sets. Some useful graphical displays include:

- mosaic displays showing the marginal relations among the response variables and of the explanatory variables, each collapsed over the other set;
- conditional mosaics or fourfold displays of the associations among the responses, stratified by one or more of the explanatory variables;

- plots of empirical logits and log odds ratios, as in Figure 8.19 or model-based plots, such as Figure 8.22, showing a model-smoothed summary.

These displays can, and should, inform our search for an adequate descriptive or explanatory model. Some of these ideas are illustrated in the following example.

{ex:toxaemia}

EXAMPLE 8.12: Toxaemic symptoms in pregnancy

Brown *et al.* (1983) gave the data used here on two signs of *toxaemia*, an abnormal condition during pregnancy characterized by high blood pressure (hypertension) and high levels of protein in the urine. If untreated, both the mother and baby are at risk of complications or death. The data frame `Toxaemia` in `vcdExtra` represents 13,384 expectant mothers in Bradford, England in their first pregnancy, who were also classified according to social class and the number of cigarettes smoked per day.

There are thus two response variables, and two explanatory variables in this data set in frequency form. For convenience, we also convert it to a $2 \times 2 \times 5 \times 3$ table.

```
data("Toxaemia", package="vcdExtra")
str(Toxaemia)

## 'data.frame': 60 obs. of 5 variables:
## $ class: Factor w/ 5 levels "1","2","3","4",...: 1 1 1 1 1 1 1 1 1 1 ...
## $ smoke: Factor w/ 3 levels "0","1-19","20+": 1 1 1 1 2 2 2 3 3 ...
## $ hyper: Factor w/ 2 levels "Low","High": 2 2 1 1 2 2 1 1 2 2 ...
## $ urea : Factor w/ 2 levels "Low","High": 2 1 2 1 2 1 2 1 2 1 ...
## $ Freq : int 28 82 21 286 5 24 5 71 1 3 ...

tox.tab <- xtabs(Freq~class + smoke + hyper + urea, Toxaemia)
ftable(tox.tab, row.vars=1)

##      smoke      0      1-19      20+
##      hyper Low High Low High Low High Low High Low High
##      urea Low High Low High Low High Low High Low High
## class
## 1      286    21    82    28    71    5    24    5    13    0    3    1
## 2      785    34   266    50   284   17    92   13    34    3   15    0
## 3     3160   164  1101   278  2300  142   492  120   383   32   92   16
## 4      656    52   213    63   649   46   129   35   163   12   40    7
## 5      245    23    78    20   321   34    74   22    65    4   14    7
```

The questions of main interest are how the occurrence of each symptom varies with social class and smoking, and how the association between these symptoms varies. It is useful, however, to examine first the marginal relationship between the two responses, and between the two predictors. The calls to `mosaic()` below produce the two panels in Figure 8.23.

```
mosaic(~smoke + class, data=tox.tab, shade=TRUE,
      main="Predictors", legend=FALSE)
mosaic(~hyper + urea, data=tox.tab, shade=TRUE,
      main="Responses", legend=FALSE)
```

We see in Figure 8.23 that the majority of the mothers are in the third social class, and that smoking is negatively related to social class, with the highest levels of smoking in classes 4 and 5. (Social class 1 is the highest in status here.) More than 50% are non-smokers. Within the responses, the great majority of women exhibit neither symptom, but showing one symptom makes it much more likely to show the other. Marginally, hypertension is somewhat more prevalent than protein urea.




Figure 8.23: Mosaic displays for Toxaemia data: Predictor and response associations fig:tox-mosaic1

We next examine how the association between responses varies with social class and with smoking. Figure 8.24 shows a collection of conditional mosaic plots using `cotabplot()` of the association between hypertension and urea, for each level of smoking, collapsed over social class.

```
cotabplot(~hyper + urea | smoke, tox.tab, shade=TRUE,
          legend=FALSE, layout=c(1, 3))
```




Figure 8.24: Toxaemia data: Response association conditioned on smoking level fig:tox-mosaic2

Figure 8.25 is similar, but stratified by social class. The marginal frequencies of the conditioning variable is not represented in these plots. (For example, as can be seen in Figure 8.23, the greatest number of women are in class 3.)

```
cotabplot(~hyper + urea | class, tox.tab, shade=TRUE,
          legend=FALSE, layout=c(1, 5))
```




Figure 8.25: Toxaemia data: Response association conditioned on social class fig:tox-mosaic3

Ignoring social class, the association between hypertension and protein urea decreases with smoking. Ignoring smoking, the association is greatest in social class 3. However, these displays don't show directly how the two symptoms are associated in the combinations of social class and smoking. The fourfold display in Figure 8.26, does that.

```
fourfold(aperm(tox.tab), fontsize=16)
```

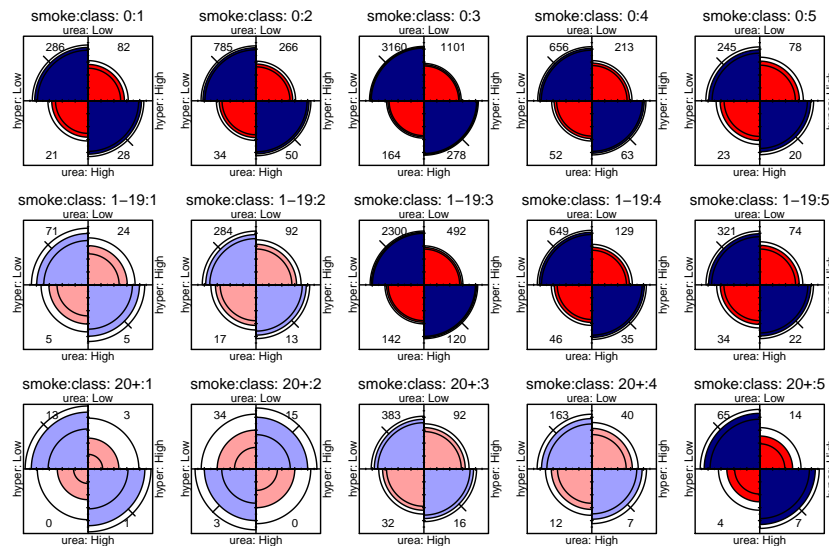


Figure 8.26: Fourfold display for Toxaemia data. Smoking levels vary in the rows and social class in the columns.

{fig:tox-fourfold}

It can be seen in Figure 8.26 that the odds ratio appears to increase with both smoking and social class number and these two symptoms are positively associated in nearly all cases. In only two cases the odds ratio is not significantly different from 1: mothers in classes 1 and 2, who smoke more than 20 cigarettes a day, but the frequency in this cell is quite small.

```
t(apply(tox.tab, MARGIN=1:2, FUN=sum))
```

```
##      class
## smoke   1   2   3   4   5
##   0  417 1135 4703 984 366
##   1-19 105  406 3054 859 451
##   20+   17   52  523 222  90
```

From these plots, it is useful to examine the association between hypertension and urea more directly, by calculating and plotting the odds ratios. For a $2 \times 2 \times K \times L \times \dots$ table, the function `oddsratio()` in `vcd` calculates these for each 2×2 subtable, and returns an array of dimension $K \times L \times \dots$, together with similar array of standard errors.

```
LOR <-oddsratio(aperm(tox.tab))
LOR
```

##		1	2	3	4	5
##	0	1.537	1.4678	1.5821	1.3168	1.005
##	1-19	1.085	0.8589	1.3739	1.3423	1.032
##	20+	2.449	-1.1458	0.7331	0.8659	2.095

The `plot()` method for the resulting "logoddsratio" object only handles a single stratum dimension, but in the present case it is easy to plot the result using `matplot()` as we did earlier. The lines below produce Figure 8.27.

```
matplot(t(LOR), type="b",
        cex=1.5, pch=15:17, cex.lab=1.5, lwd=2, lty=1,
        ylab='log odds ratio: Urea | Hypertension',
        xlab='Social class of mother',
        xlim=c(1, 5.5), col=c("blue", "black", "red"))
)
abline(h=0, col='gray')
text(5.2, LOR[,5]+c(-.05,.05), labels=rownames(LOR), cex=1.25)
text(5.2, max(LOR[,5])+.2, "Smoking", cex=1.4)
```




Figure 8.27: Log odds ratios for protein urea given hypertension, by social class and level of maternal smoking

The association between the response symptoms, shown in Figure 8.27 is clearer, once we take the variation in sample sizes into account. Except for the heavy smokers, particularly in social classes 1 and 2, the log odds ratio appears to range only between 1–1.5, meaning that, given one symptom, the odds of also having the other range between $\exp(1) = 2.72$ and $\exp(1.5) = 4.48$.

This initial overview of the data is completed by calculating and plotting the log odds for each symptom within each class-smoke population. This could be done in the same way as in Example 8.11, (except that there are now two explanatory factors). The steps used there were: (a) Reshape the $2 \times 2 \times K \dots$ table to a matrix with four columns corresponding to the binary response combinations. (b) Calculate the logits (and log odds ratio) using `blogits()`. **TODO:** Use this as an exercise.

Here, it is more useful to make separate plots for each of the logits, and we illustrate a more general approach that applies to two or more binary responses, with two or more predictor variables. The essential idea is to fit a separate logit model for each response separately, using the *highest-order interaction* of all predictors (the saturated model). The fitted logits in these models then match those in the data.

```
tox.hyper <- glm(hyper=='High' ~ class*smoke, weights=Freq,
                 data=Toxaemia, family=binomial)
tox.urea <- glm(urea=='High' ~ class*smoke, weights=Freq,
                 data=Toxaemia, family=binomial)
```

It is then simple to plot these results using the `effects` package as shown in Figure 8.28. Each plot shows the logit for the response measure against class, with separate curves for the levels of smoking.¹⁵

```
library(effects)

plot(allEffects(tox.hyper),
     ylab = "Probability (hypertension)",
     xlab = "Social class of mother",
     main = "Hypertension: class*smoke effect plot",
     colors = c("blue", "black", "red"),
     lwd=3, multiline=TRUE,
     key.args=list(x=0.05, y=0.2, cex=1.2)
)

plot(allEffects(tox.urea),
     ylab = "Probability (Urea)",
     xlab = "Social class of mother",
     main = "Urea: class*smoke effect plot",
     colors = c("blue", "black", "red"),
     lwd=3, multiline=TRUE,
     key.args=list(x=0.65, y=0.2, cex=1.2)
)
```

From Figure 8.28, it can be seen that the prevalence of these symptoms has a possibly complex relation to social class and smoking. However, the mosaic for these predictors in Figure 8.23 has shown us that several of the class-smoking categories are quite small (particularly heavy smokers in Classes 1 and 2) so the response effects for these classes will be poorly estimated. Taking this into account, we suspect that protein urea varies with social class, but not with smoking, while the prevalence of hypertension may truly vary with neither, just one, or both of these predictors.

Fitting models

The plots shown so far in this example are all essentially *data-based*, in that they use the observed frequencies or transformations of them and don't allow for a simpler view, based on a reasonable

¹⁵As is usual for effect plots of binary response `glm()` models, the vertical axis is plotted on the scale of log odds, but labeled in terms of probabilities.




Figure 8.28: Plots of log odds for hypertension and urea, by social class of mother and smoking. fig:tox-effpl

model. That is, abbreviating the table variables by their initial letters, the plots in Figure 8.27 and Figure 8.28 are plots of the saturated model, [CSHU] that fits perfectly, but with the data transformed for each 2×2 subtable to the log odds ratio and the two log odds for `hyper` and `urea`.

The bivariate logistic model fit by `vglm()` still applies when there are two or more predictors; however, like other multivariate response models, it doesn't easily allow the logits to depend on *different* predictor terms. To illustrate this, we first transform the `Toxaemia` to a 15×4 data frame in the form required by `vglm()`.

```
tox.tab <- xtabs(Freq~class + smoke + hyper + urea, Toxaemia)
toxaemia <- t(matrix(aperm(tox.tab), 4, 15))
colnames(toxaemia) <- c("hu", "hU", "Hu", "HU")
rowlabs <- expand.grid(smoke=c("0", "1-19", "20+"), class=factor(1:5))
toxaemia <- cbind(toxaemia, rowlabs)
head(toxaemia)
```

```
##      hu hU  Hu HU smoke class
## 1 286 21  82 28      0      1
## 2  71  5  24  5    1-19      1
## 3  13  0   3  1     20+      1
## 4 785 34 266 50      0      2
## 5 284 17  92 13    1-19      2
## 6  34  3  15  0     20+      2
```

In the model specification for `vglm()`, the `zero` argument in `binom.or()` allows any one or more of the two log odds and log odds ratio to be fit as a constant (intercept-only) in Eqn. (8.29). However, in that equation, the predictors x_1, x_2, x_{12} , must be the *same* in all three submodels. For example, the model `tox.vglm1` below uses main effects of `class` and `smoke` in both models for the logits, and `zero=3` for a constant log odds ratio.

```
tox.vglm1 <- vglm(cbind(hu, hU, Hu, HU) ~ class + smoke,
                 binom2.or(zero=3), data=toxaemia)
coef(tox.vglm1, matrix=TRUE)
```

```
##          logit(mu1) logit(mu2) log(oratio)
## (Intercept) -0.5085365 -1.221452      2.781
## class2      0.1815646  0.038205      0.000
## class3      0.0633277 -0.008755      0.000
## class4     -0.0222706 -0.003154      0.000
## class5     -0.0007717  0.082186      0.000
## smoke1-19  -0.4129865 -0.219867      0.000
## smoke20+   -0.3056247 -0.124502      0.000
```

Instead, when there are no quantitative predictors, and when the odds ratio is relatively constant (as here) it is easier to fit ordinary loglinear models than to use the bivariate logit formulation of the previous example. These allow the responses H and U to depend on the class-smoking combinations separately, by including the terms $[CSH]$ or $[CSU]$, respectively.

The minimal, null model, $[CS][H][U]$ fits the marginal association of the numbers in each class-smoking category, but asserts that the responses, H and U are independent, which we have already seen is contradicted by the data. We take $[CS][HU]$ as the baseline model (Model 1), asserting no relation between response and predictor variables, but associations within each set are allowed. These models are fit as shown below.

```
# null model
tox.glm0 <- glm(Freq ~ class*smoke + hyper + urea,
               data=Toxaemia, family=poisson)
# baseline model: no association between predictors and responses
tox.glm1 <- glm(Freq ~ class*smoke + hyper*urea,
               data=Toxaemia, family=poisson)
```

We proceed to fit a collection of other models, adding terms to allow more associations between the responses and predictors. Summary measures of goodness of fit and parsimony are shown in Table 8.4.

Table 8.4: Loglinear models, `tox.glm*`, fit to the Toxaemia data

{tab:toxmod}

Model	Terms	df	G^2	p -value	G^2/df	AIC	BIC	R^2
0	CS H U	43	672.85	0.0000	15.65	586.85	264.27	.
1	CS HU	42	179.03	0.0000	4.26	95.03	-220.04	0.000
2	CS HU SH CU	36	46.12	0.1203	1.28	-25.88	-295.94	0.742
3	CS CH CU HU SH CU	30	40.47	0.0960	1.35	-19.53	-244.58	0.774
4	CSH CU HU	24	26.00	0.3529	1.08	-22.00	-202.04	0.855
5	CSH CU SU HU	22	25.84	0.2588	1.17	-18.16	-183.20	0.856
6	CSH CSU HU	14	22.29	0.0729	1.59	-5.71	-110.74	0.875
7	CSH CSU SHU	12	15.65	0.2079	1.30	-8.35	-98.37	0.913
8	CSH CSU CHU SHU	8	12.68	0.1233	1.59	-3.32	-63.33	0.929
9	CSHU	0	0.00	0	0	0.00	0.00	1.000

```
tox.glm2 <- update(tox.glm1, . ~ . + smoke*hyper + class*urea)
tox.glm3 <- glm(Freq ~ (class + smoke + hyper + urea)^2,
               data=Toxaemia, family=poisson)
tox.glm4 <- glm(Freq ~ class*smoke*hyper + hyper*urea + class*urea,
```



```

      data=Toxaemia, family=poisson)
tox.glm5 <- update(tox.glm4, . ~ . + smoke*urea)
tox.glm6 <- update(tox.glm4, . ~ . + class*smoke*urea)
tox.glm7 <- update(tox.glm6, . ~ . + smoke*hyper*urea)
tox.glm8 <- glm(Freq ~ (class + smoke + hyper + urea)^3,
  data=Toxaemia, family=poisson)
tox.glm9 <- glm(Freq ~ (class + smoke + hyper + urea)^4,
  data=Toxaemia, family=poisson)

```

Model 2 adds the simple dependence of hypertension on smoking ($[SH]$) and that of urea on class ($[CU]$). Model 3 includes all two-way terms. In Model 4, hypertension is allowed to depend on both class and smoking jointly ($[CSH]$). In Model 5 an additional dependence of urea on smoking ($[SU]$) is included, while in Model 6 urea depends on class and smoking jointly ($[CSU]$).

None of these models contain three-way terms involving both H and U , so these models assume that the log odds ratio for hypertension given urea is constant over the explanatory variables. Recalling the conditional mosaics (Figure 8.24 and Figure 8.25), Models 7 and 8 add terms which allow the odds ratio to vary, first with smoking ($[SHU]$), then with class ($[CHU]$) as well. Finally, Model 9 is the saturated model, that fits perfectly.

How do we choose among these models? Model 2 is the smallest model whose deviance is non-significant. Models 4 and 5 both have a smaller ratio of G^2/df . For comparing nested models, we can also examine the change in deviance as terms are added (or dropped). Thus, going from Model 2 to Model 3 decreases the deviance by 5.65 on 6 df, while the step from Model 3 to Model 4 gives a decrease of 14.47, also on 6 df. These tests can be performed using `lrtest()` in the `lmtest` package, shown below for models `tox.glm1`–`tox.glm5`.

```

library(lmtest)
lrtest(tox.glm1, tox.glm2, tox.glm3, tox.glm4, tox.glm5)

## Error: unused arguments (tox.glm3, tox.glm4, tox.glm5)

```

The AIC and BIC statistics, balancing parsimony and goodness-of-fit, have their minimum value for Model 2, which we adopt here for this example.

Plotting model results

Whatever model is chosen, as a final step, it is important to determine what that model implies about the original research questions. Because our focus here is on the prevalence of each symptom, and their association, it is helpful to graph the fitted logits and log odds ratios implied by the model, as was done in Figure 8.21 and Figure 8.22.

The presentation goal here is to produce plots showing the observed logits and log odds ratios as in Figure 8.28 and Figure 8.27, supplemented by lines showing these values according to the fitted model. In Example 8.11 we fit the bivariate logit model, for which the response functions were the desired logits and log odds. Here, where we have fit ordinary loglinear models, the observed and fitted logits can be calculated from the observed and fitted frequencies. The calculations require a bit of R calisthenics to arrange these into forms suitable for plotting.

As we did earlier, we first reshape the `Toxaemia` to wide format, as a 15×4 table of observed frequencies. Because there are now two predictor variables, we take care to include the levels of `smoke` and `class` as additional columns.

```
# reshape to 15 x 4 table of frequencies
tox.tab <- xtabs(Freq~class + smoke + hyper + urea, Toxaemia)
toxaemia <- t(matrix(aperm(tox.tab), 4, 15))
colnames(toxaemia) <- c("hu", "hU", "Hu", "HU")
rowlabs <- expand.grid(smoke=c("0", "1-19", "20+"), class=factor(1:5))
toxaemia <- cbind(toxaemia, rowlabs)
```

Applying `blogits()`, we get the observed logits and log odds ratios in `logitsTox`.

```
# observed logits and log odds ratios
logitsTox <- blogits(toxaemia[,4:1], add=0.5)
colnames(logitsTox)[1:2] <- c("logitH", "logitU")
logitsTox <- cbind(logitsTox, rowlabs)
head(logitsTox)

##      logitH logitU logOR smoke class
## 1 -1.0206 -1.999  1.527    0     1
## 2 -0.9426 -2.166  1.071  1-19    1
## 3 -1.0296 -2.140  2.449   20+    1
## 4 -0.9504 -2.516  1.462    0     2
## 5 -1.0470 -2.498  0.864  1-19    2
## 6 -0.8650 -2.526 -1.146   20+    2
```

The fitted frequencies are extracted using `predict(tox.glm2, type="response")` and then manipulated in a similar way to give `logitsFit`.

```
# fitted frequencies, as a 15 x 4 table
Fit <- t(matrix(predict(tox.glm2, type="response"), 4, 15))
colnames(Fit) <- c("HU", "Hu", "hU", "hu")
Fit <- cbind(Fit, rowlabs)
logitsFit <- blogits(Fit[,1:4], add=0.5)
colnames(logitsFit)[1:2] <- c("logitH", "logitU")
logitsFit <- cbind(logitsFit, rowlabs)
```

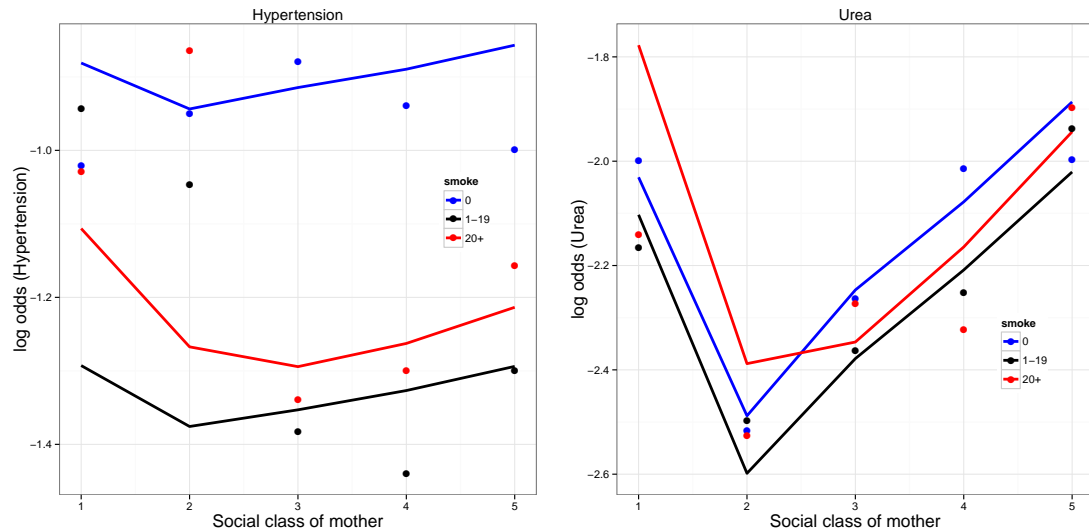
In tabular form, you can examine any of these components, for example, the log odds ratios from the fitted values shown below.

```
matrix(logitsFit$logOR, 3, 5,
       dimnames=list(smoke=c("0", "1-19", "20+"), class=1:5))


##      class
## smoke    1      2      3      4      5
## 0      1.359 1.364 1.368 1.364 1.358
## 1-19 1.358 1.368 1.368 1.367 1.366
## 20+  1.280 1.347 1.366 1.362 1.351
```

Finally, we can plot the observed values in `logitsTox` (as points) and the fitted values under Model 2 in `logitsFit` (as lines), separately for the `logitH`, `logitU`, and `logOR` components. The code below uses `ggplot2` for the log odds of hypertension, and is repeated for urea and the log odds ratio. These graphs are shown in Figure 8.29 and Figure 8.30.

```
ggplot(logitsFit, aes(x=as.numeric(class), y=logitH, color=smoke)) +
  theme_bw() +
  geom_line(size=1.2) +
  scale_color_manual(values=c("blue", "black", "red")) +
  ylab("log odds (Hypertension)") +
  xlab("Social class of mother") +
```

Figure 8.29: Observed (points) and fitted (lines) logits for the `Toxemia` data under Model 2.

{fig:tox-glm-logits3}

Figure 8.30: Observed (points) and fitted (lines) log odds ratios for the `Toxemia` data under Model 2.

{fig:tox-glm-logits3}

```
ggtitle("Hypertension") +
  theme(axis.title=element_text(size=16)) +
  geom_point(data=logitsTox,
    aes(x=as.numeric(class), y=logitH, color=smoke), size=3) +
  theme(legend.position=c(0.85, .6))
```

According to this model, Figure 8.30 shows that the fitted log odds ratio is in fact nearly constant, while Figure 8.29 shows that the log odds for hypertension depends mainly on smoking (with a large difference of the non-smoking mothers from the rest) and that for protein urea depends mainly on social class.¹⁶

¹⁶Some possible enhancements to these graphs include (a) plotting on the scale of probabilities or including a right

Yet, the great variability of the observed points around the fitted curves indicates that these relationships are not well-determined. Adding error bars showing the standard error around each fitted point would indicate that the data conforms as closely to the model as can be expected, given the widely different sample sizes. However, this would make the plots more complex, and so was omitted here. In addition to showing the pattern of the results according to the fitted model, such graphs also help us to appreciate the model's limitations.

△

8.11 Chapter summary

{sec:loglin-summary}

- Loglinear models provide a comprehensive scheme to describe and understand the associations among two or more categorical variables. It is helpful to think of these as discrete analogs of ANOVA models, or of regression models, where the log of cell frequency is modelled as a linear function of predictors.
- Loglinear models typically make no distinction between response and explanatory variables. When one variable *is* a response, however, any logit model for that response has an equivalent loglinear model. The logit form is usually simpler to formulate and test, and plots of the observed and fitted logits are easier to interpret.
- Models for square tables, with the same row and column categories are an important special case. For these and other structured tables, a variety of techniques provide the opportunity to fit models more descriptive than the independence model and more parsimonious than the saturated model.
- Standard loglinear models treat all variables as unordered factors. When one or more factors are ordinal, however, loglinear and logit models may be simplified by assigning quantitative scores to the levels of an ordered factor. Such models are often more sensitive and have greater power because they are more focused.
- In all these cases, the interplay between graphing and fitting is important in arriving at an understanding of the relationships among variables and an adequate descriptive model which is faithful to the details of the data.
- When there are several categorical responses, along with one or more explanatory variables, some special forms of loglinear and logit models may be used to separate the marginal dependence of each response on the explanatory variables from the interdependence among the responses.

8.12 Further reading

{sec:loglin-reading}

8.13 Lab exercises

{sec:loglin-lab}

1. Example 8.7 presented an analysis of the data on visual acuity for the subset of women in the `VisualAcuity` data. Carry out a parallel analysis of the models fit there for the men in this data set, given by:

vertical axis showing corresponding probabilities; (b) using the same vertical axis limits for the two graphs for direct comparison.

```
data("VisualAcuity", package="vcd")
men <- subset(VisualAcuity, gender=="male", select=-gender)
```

2. Table 8.5 gives a 4×4 table of opinions about premarital sex and whether methods of birth control should be made available to teenagers aged 14–16 from the 1991 General Social Survey (Agresti, 2013, Table 10.3). Both variables are ordinal, and their grades are represented by the case of the row and column labels.

Table 8.5: Opinions about premarital sex and availability of teenage birth control. *Source:* Agresti (2013, Table 10.3)

{tab:birthcontrol}

Premarital sex	Birth control			
	DISAGREE	disagree	agree	AGREE
WRONG	81	68	60	38
Wrong	24	26	29	14
wrong	18	41	74	42
OK	36	57	161	157

- Fit the independence model to these data using `loglm()` or `glm()`.
- Make a mosaic display showing departure from independence and describe verbally the pattern of association.
- Treating the categories as equally spaced, fit the $L \times L$ model of uniform association, as in Section 8.6. Test the difference against the independence model with a likelihood-ratio test.
- Fit the RC(1) model with `gnm()`, and test the difference of this against the model of uniform association.
- Write a brief summary of these results, including plots useful for explaining the relationships in this data set.

References

- Agresti, A. (1984). *Analysis of Ordinal Categorical Data*. New York: Wiley.
- Agresti, A. (1990). *Categorical Data Analysis*. New York: Wiley-Interscience.
- Agresti, A. (2013). *Categorical Data Analysis*. Wiley Series in Probability and Statistics. New York: Wiley-Interscience [John Wiley & Sons], 3rd edn.
- Akaike, H. (1973). Information theory and an extension of the maximum likelihood principal. In B. N. Petrov and F. Czaki, eds., *Proceedings of the 2nd International Symposium on Information*. Budapest: Akademiai Kiado.
- Ashford, J. R. and Sowden, R. D. (1970). Multivariate probit analysis. *Biometrics*, 26, 535–546.
- Birch, M. W. (1963). Maximum likelihood in three-way contingency tables. *Journal of the Royal Statistical Society, Series B*, 25, 220–233.
- Bishop, Y. M. M., Fienberg, S. E., and Holland, P. W. (1975). *Discrete Multivariate Analysis: Theory and Practice*. Cambridge, MA: MIT Press.
- Brown, P. J., Stone, J., and Ord-Smith, C. (1983). Toxaemic signs during pregnancy. *Journal of the Royal Statistical Society, Series C (Applied Statistics)*, 32, 69–72.
- Brunswick, A. F. (1971). Adolescent health, sex, and fertility. *American Journal of Public Health*, 61(4), 711–729.
- Caussinus, H. (1966). Contribution à l'analyse statistique des tableaux de corrélation. *Annales de la Faculté des Sciences de l'Université de Toulouse*, 39 (année 1965), 77–183.
- Evers, M. and Namboodiri, N. K. (1977). A Monte Carlo assessment of the stability of log-linear estimates in small samples. In *Proceedings of the Social Statistics Section*. Alexandria, VA: American Statistical Association.
- Fienberg, S. E. (1980). *The Analysis of Cross-Classified Categorical Data*. Cambridge, MA: MIT Press, 2nd edn.
- Fienberg, S. E. and Rinaldo, A. (2007). Three centuries of categorical data analysis: Log-linear models and maximum likelihood estimation. *Journal of Statistical Planning and Inference*, 137(11), 3430–3445.
- Firth, D. (2003). Overcoming the reference category problem in the presentation of statistical models. *Sociological Methodology*, 33, 1–18.
- Firth, D. and Menezes, R. X. d. (2004). Quasi-variances. *Biometrika*, 91, 65–80.

- Goodman, L. A. (1970). The multivariate analysis of qualitative data: Interactions among multiple classifications. *Journal of the American Statistical Association*, 65, 226–256.
- Goodman, L. A. (1971). The analysis of multidimensional contingency tables: Stepwise procedures and direct estimates for building models for multiple classifications. *Technometrics*, 13, 33–61.
- Goodman, L. A. (1972). Some multiplicative models for the analysis of cross classified data. In *Proceedings of the sixth Berkeley Symposium on Mathematical Statistics and Probability*, (pp. 649–696). Berkeley, CA: University of California.
- Goodman, L. A. (1979). Simple models for the analysis of association in cross-classifications having ordered categories. *Journal of the American Statistical Association*, 74, 537–552.
- Goodman, L. A. (1985). The analysis of cross-classified data having ordered and/or unordered categories: Association models, correlation models, and asymmetry models for contingency tables with or without missing entries. *Annals of Statistics*, 13(1), 10–69.
- Goodman, L. A. (1986). Some useful extensions of the usual correspondence analysis approach and the usual log-linear models approach in the analysis of contingency tables. *International Statistical Review*, 54(3), 243–309. With a discussion and reply by the author.
- Haberman, S. J. (1972). Statistical algorithms: Algorithm AS 51: Log-linear fit for contingency tables. *Applied Statistics*, 21(2), 218–225.
- Haberman, S. J. (1973). The analysis of residuals in cross-classified tables. *Biometrics*, 29, 205–220.
- Haberman, S. J. (1974). *The Analysis of Frequency Data*. Chicago: University of Chicago Press.
- Hauser, R. M. (1979). Some exploratory methods for modeling mobility tables and other cross-classified data. In K. F. Schuessler, ed., *Sociological Methodology 1980*, (pp. 413–458). San Francisco: Jossey-Bass.
- Lang, J. B. and Agresti, A. (1994). Simultaneously modeling joint and marginal distributions of multivariate categorical responses. *Journal of the American Statistical Association*, 89(426), 625–632.
- Lindsey, J. K. (1995). *Modelling Frequency and Count Data*. Oxford, UK: Oxford University Press.
- McCullagh, P. and Nelder, J. A. (1989). *Generalized Linear Models*. London: Chapman and Hall.
- Powers, D. A. and Xie, Y. (2008). *Statistical Methods for Categorical Data Analysis*. Bingley, UK: Emerald, 2nd edn.
- Raftery, A. E. (1986). Choosing models for cross-classifications. *American Sociological Review*, 51, 146–146.
- Schwartz, G. (1978). Estimating the dimensions of a model. *Annals of Statistics*, 6, 461–464.